
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 A printed claim that the scorecaster breaks online conformal coverage meets its
2 refutation by overlapping authors; this paper joins them and locates the edge
3 of the set that refutation relies on. Conformal PID states long-run coverage for
4 any bounded scorecaster; a corollary keeps it for predictable perturbations of the
5 deployed threshold inside an admissible radius. The claim was acted on; placement
6 and derivation are conceded by name. At the null scorecaster that condition is *tight*.
7 A legal L_1 dead band on the *completed* threshold leaves that set at its radius, and
8 past it miscoverage is 1.000000 against $\alpha = 0.1$. Coverage goes, not only the rate.
9 The edge belongs to the saturator–scorecaster pair, not the readout; elsewhere the
10 condition is sufficient only. Under the equally legal $\hat{q} \equiv +b/2$ the failing band
11 covers; at $\hat{q} \equiv -b/2$ both it and partial adjustment at $w = 0.999$ miscover.

12 1 Introduction

13 An online conformal method emits a threshold before each observation and moves it. Forecast
14 stability names, scores and prices it [Tunc et al., 2013, Godahewa et al., 2025, Van Belle et al.,
15 2023, Pritularga and Kourentzes, 2024, Genov et al., 2026, Van Belle et al., 2026]. Varying temporal
16 structure at fixed level came earlier, in forecast verification, matching predictive marginals on real
17 producers [Gneiting et al., 2007, Pinson and Girard, 2012, Worsnop et al., 2018]. Coverage and mean
18 length under-describe a deployed interval [Min et al., 2026, Vaze, 2026]. None of it is claimed here.

19 **First: the record carries a claim and its refutation, and this paper joins them.** Conformal
20 PID [Angelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator, asking only
21 that it be predictable and lie in $[-b/2, b/2]$: “any function of the past” (Theorem 1). A readout
22 there is already quantified over. Also in print, of that method: “[s]ince using a scorecaster (D-part)
23 in C-PID breaks the theoretical coverage guarantee” [Li and Rodríguez, 2025], stated and acted
24 on in its baselines paragraph, verbatim on refereed printed page 18443. It is refuted twice, with
25 proof. Li et al. [2025, Corollary 2] add an arbitrary predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ to the
26 deployed threshold, note that the result “has exactly the CPID error-integration form”, and conclude
27 $|E_T| = o(T)$. The same generic-slot argument carries AcMCP’s long-run coverage [Wang and
28 Hyndman, 2024]. The first and last authors of Li et al. [2025] wrote Li and Rodríguez [2025] and cite
29 it: the halves already share a bibliography, this is ordinary self-correction, and the joining is what
30 this paper adds. Four surfaces were queried on 20 August 2026 for both identifiers and “scorecaster”.
31 They are arXiv full text, arXiv abstracts, OpenReview, and the six works Semantic Scholar records
32 citing Li and Rodríguez [2025]. None returns both the claim and the corollary. Hu et al. [2026] still
33 call scorecaster choice “arbitrary” and lacking “principled guidance” while asserting valid coverage:
34 model selection, not validity. **We claim neither the placement nor the derivation.**

35 **Second: that admissible set has an edge, and the condition there is tight.** Corollary 2 is a
36 *retention* result: what it leaves open is outside. Section 3 exhibits a legal perturbation leaving it, the
37 L_1 dead-band family characterised here on the *completed* threshold, whose edge is the corollary’s




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In (B) the loop closes on the integrator’s own output, and both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set. In (A) it closes on $\tilde{q}_{t+1}$, which it never computed, and the deployed sequence leaves that set.

38 radius at the null scorecaster. Past it coverage goes, not only the rate: the published sufficient condition
39 is necessary there. Section 4 sets both claims beside four vocabularies for the same movement.

40 2 Setup

41 **The producer.** A forecaster fit before t induces a score s_t , and a threshold q_t is emitted before
42 y_t is revealed. The deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = \mathbf{1}\{s_t(x_t, y_t) > q_t\}$. The
43 target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID
44 [Angelopoulos et al., 2023] manipulates the threshold on the score scale,

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

45 their iteration (5). Its r_t saturates: condition (4) requires $|r_t(x)| \geq b$ with the sign of x once $|x| \geq$
46 $c h(t)$, for h nonnegative, nondecreasing and sublinear. With $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$, Proposition 2
47 gives $|E_T| \leq c h(T) + 1$, read on realised scores.

48 **One scalar, two readouts, two places.** Damping a sequentially updated scalar has two standard
49 readouts, both prior work. A quadratic cost gives linear partial adjustment $\tilde{q}_{t+1} = (1 - \lambda)\hat{q}_{t+1}^{\text{raw}} + \lambda\tilde{q}_t$
50 [Gârleanu and Pedersen, 2013], published as post-processing [Godahewa et al., 2025] and applied
51 to a deployed threshold [Binny and Dixit, 2025, Eq. 13]. A proportional cost gives a no-trade band,
52 moving only when the target is more than τ away, then pulling back exactly τ [Constantinides,
53 1986, Davis and Norman, 1990]. Either readout has two positions in (1), drawn in Figure 1; (A) is
54 measured here. *The dead band (L_1) is not a stylistic alternative to the partial-adjustment (L_2) form;*
55 *it is the family whose failure Section 3 characterises. Neither is safe by construction: L_2 forfeits*
56 *Proposition 2’s rate while covering, L_1 past τ^* loses coverage, and at the legal $\hat{q} \equiv -b/2$ both lose it.*
57 *What is tight is the radius, not a form.*

58 **The quantity, and the names it already has.** We report the deployed threshold’s *travel*, $\mathcal{T}_H =$
59 $\sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new. Over a predictive
60 distribution $\sum |\Delta q|$ is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls it
61 *jumpiness*, a *convergence index*, (in)consistency [Zsótér et al., 2009, Ehret, 2010, Pappenberger
62 et al., 2011]; applied control, IACER [Sarhadi, 2025]; dynamic-regret online learning, a path length.
63 Table 2’s caption names ten, so a fresh one is used here.

64 **Harness.** The harness fixes $\alpha = 0.1$ and $b = 2$. Its saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$,
65 with $c = 1$ and $h(t) = \log(t + 2)$, meets (4) with equality. The scorecaster is null and legal, reducing
66 (1) to the pure error integration Proposition 2 addresses. Neither choice is neutral: (4) is a lower
67 bound met exactly here, and Theorem 1 permits scorecasters $b/2$ either side of null. The adversary is
68 specified, not tie-broken. It plays $s_t = \text{clip}(\hat{q}_t + \varepsilon, \pm b/2)$, $\varepsilon = 10^{-9}$, against the *deployed* $\hat{q}_t$: the
69 $\varepsilon \downarrow 0$ reading of “sets the score at the threshold”, firing the strict indicator. Legality is Theorem 1’s
70 $[-b/2, b/2]$ on both sides, so $b/2 + b/2 = b$ is exactly tight.

71 3 The edge of the admissible set, measured

72 **The retention result is published, and tested from outside.** Li et al. [2025, Corollary 2] prove
73 that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the deployed threshold keeps (1) in error-integration
74 form and *retains* $|E_T| \leq c h(T) + 1$: the radius is theirs. A downstream partial adjustment of weight

Table 1: Placement A: a one-scalar readout on the *completed* threshold; one deterministic adversary, one pass to $T = 10^6$. *All eleven arms run are listed*, so no covering arm is withheld; five exceed the bound at $T = 10^6$. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, “Sat.” the fraction of rounds on which (4) delivers its full $\pm b$. Every figure is the *primary regime*’s, defined in Appendix A, which also gives the $T = 10^4$ excursions and the fitted excursion law.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5




Figure 2: The dead-band coverage cliff, and the fact that its edge moves. **(a)** Realised miscoverage against dead-band width τ under five admissible $(r_t, \hat{q})$ pairs. Each pair’s $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$ is marked in its own colour, dashed where it falls on the axis and arrowed where it does not: 1 at the null scorecaster, 2 at $\hat{q} \equiv +b/2$ and 7 at a saturator of level $4b$ are dashed; 0 at $\hat{q} \equiv -b/2$ and unbounded under Angelopoulos et al.’s tangent integrator are arrows. Three widths run under each pair, plus four wider bands at the null scorecaster. All 19 dead-band runs land on the side of their own τ^* that the law predicts, and the vertical dashes are the law’s prediction, not four located edges: only the null scorecaster’s edge has runs on both sides of τ^* . Appendix A gives the realised miscoverage at each covering setting. **(b)** An eleven-point grid at the null scorecaster *locates* that edge between $\tau = 1.000$ and 1.001 , with no intermediate regime. Horizons: $T = 10^6$ in (a) except the four widest bands ($\tau \geq 2$), run only to $T = 2 \times 10^5$; $T = 10^5$ in (b).

75 w leaves it; Table 1 prices that departure. For $w = 0.99$ and 0.999 the excursion is *flat* in T while the
76 bound grows like $\log T$, so the ratio falls with the horizon. At $w = 0.999$ it is 61.08 at $T = 10^4$ and
77 42.10 at $T = 10^6$. The forfeit is a property of the smoothing *strength*, not Placement A, and it buys
78 something: travel falls from 28,311 to 202. The forfeit’s sign is counter-intuitive. A readout on the
79 completed q_{t+1} touches neither r_t nor the integrator’s argument, so it cannot put (4) at risk. It delays
80 a correction already computed, and the accumulator excurses *further*. Smoothing the output does not
81 damp E_t ; it feeds it. That 623.70 turns on the adversary’s ε , not the counting convention.

82 **The L_1 dead band leaves the radius, and the edge is measurable.** A dead band pulls the deployed
83 threshold back by exactly τ when it releases, so “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under
84 saturation the deployed value is pinned τ below the integrator’s ceiling, and its realised supremum
85 is $b - \tau$ to the printed digit at every $\tau \geq 0.9$. The adversary’s best legal play is $b/2$, so coverage is
86 lost when that supremum drops below it, at $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$. All 19 dead-band runs
87 agree with that law (Figure 2); only the null scorecaster’s edge is located, the others one-sided. **Here**
88 $\sup_x r_t(x) = b$ and $\hat{q} \equiv 0$, so $\tau^* = b/2$, one point of the admissible class, and both terms move it.
89 The legal $\hat{q} \equiv +b/2$ carries the $\tau = 1.5$ band inside the bound. The equally legal $\hat{q} \equiv -b/2$ drops the
90 edge to 0, and there $\tau = 0.5$ returns 1.000000, as does partial adjustment at $w = 0.999$. Condition (4)

91 bounds $\sup_x r_t(x)$ only from below, so raising it costs the theorem nothing; it moves the edge too and
92 shrinks the excursion (Figure 2). Under the unbounded tangent integrator Angelopoulos et al. [2023]
93 report using in all their experiments, τ^* is unbounded and all three widths there cover. The excursion
94 is 20.10 rather than 623.70, both at $T = 10^6$. **The method’s own integrator therefore sits inside**
95 **the set at every width run here.** Outside the set the transition is a step. Past τ^* the threshold sits
96 permanently below the reachable score range and $|E_t| = (1 - \alpha)t$ grows linearly, 60,747 times the
97 bound at $T = 10^6$. The wider bands $\tau = 2.0, 2.5, 3.0$ and 5.0 fail too at $T = 2 \times 10^5$, so the failing
98 set is (τ^*, ∞) , open at the left. **The endpoint is where the adversary’s specification decides the**
99 **answer.** At $\tau = \tau^*$ the deployed supremum equals the adversary’s own cap $b/2$, so the clip forces an
100 exact play and the reading settles it: the primary regime covers, while scoring that play strictly returns
101 0.000000 and scoring it as a miscoverage returns 1.000000. Appendix A separates the two readings.
102 **The failing arms have** $\text{Sat.} = 1.0000$. Condition (4) delivers its full $\pm b$ on every round, and what
103 fails is Proposition 2’s other hypothesis: the threshold closing the loop *is* the integrator’s output.

104 **What that measures is that the published condition is tight.** At $\mu_t = 1$, $\hat{q} \equiv 0$ the measured
105 τ^* is exactly Corollary 2’s radius, and crossing it costs coverage rather than the rate. **The claim is**
106 **that necessity: a tightness result about someone else’s sufficient condition**, smaller and more
107 precise than a forfeit of the rate by post-processing. Elsewhere it stays sufficient only: at $\hat{q} \equiv +b/2$
108 Corollary 2 permits radius 0 while every width run there covers. That asymmetry makes “tight” a
109 located claim.

110 **The constructive reading, and it is the corollary’s rather than ours.** Placement B keeps $\{\hat{q}_t\}$
111 inside the admissible radius, so Theorem 1 carries it already. What the measurement adds is that the
112 band has an edge, computable before deployment from $\sup_x r_t(x)$ and $\sup_t \hat{q}_t$, and that crossing it is
113 not a graceful degradation.

114 4 Where this sits

115 Four vocabularies name one quantity many ways, and the object that settles its validity (the admissible
116 set around the deployed threshold) sits in the first of them. Table 2 in Appendix B names them and
117 what each contributes. **What is conceded, by name.** The placement is prior work: Angelopoulos et al.
118 [2023] state coverage for any admissible integrator and any scorecaster, and run a smoothing-family
119 model in that slot themselves. Dupuy et al. [2026, Appendix A] give the generic argument, and
120 Duerst et al. [2024] adopt that scorecaster. The derivation is prior work too, and this concession
121 matters most. Li et al. [2025, Corollary 2] prove the retention statement in Conformal PID’s own
122 notation, with Wang and Hyndman [2024] an independent second instance. Section 2 concedes the
123 smoother object, the metric arithmetic and both readout maps. Zheng et al. [2025], Wu et al. [2025],
124 Chen et al. [2023], Lade et al. [2026] each place a smoother inside their own validity argument. One
125 vocabulary out, this is integrator anti-windup with a feedforward path: free-until-the-bound-binds
126 is its defining specification, not a finding [Galeani et al., 2009]. Nearest in the other direction, Gao
127 et al. [2025] run the same recursion but gate the *model* update on the conformal score. Their one
128 non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$ for a sublinear h , supported with a plot
129 rather than derived.

130 5 Limitations

131 **The evidence is one construction, one adversary, one saturator;** and $0.63/(1 - w)$ is specific to
132 the saturator’s *level*, which condition (4) bounds only from below. **The inherited guarantee is weak**
133 **where it applies.** Under the tangent integrator $h(t) = t/\log t$, the certified gap $(ch(T) + 1)/T$
134 decays as $O(1/\log T)$ rather than $O(1/T)$. What is forfeited was not a tight certificate. **The dead-**
135 **band widths are absolute, and the failure is total only against the adversary.** τ is in score units
136 with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$. Under i.i.d. scores it returns 0.249376 at $T = 10^6$, an
137 adversary-free confirmation of the mechanism rather than a retreat. The threshold pins at $b - \tau = 0.5$,
138 and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly. **The harness was rebuilt**
139 **against numbers already read**, which carries less weight than a rerun done without sight of them.
140 **Placement B is conceded, not tested. One seed, and only the null scorecaster’s edge is bracketed;**
141 the others are one-sided. **What would overturn the claim** is a legal $(r_t, \hat{q})$ pair with $\tau \leq \tau^*$ that
142 loses coverage, or $\tau > \tau^*$ that keeps it; the sweep has neither.

References

- Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages 741–763, 2013.
- Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. doi: 10.52202/075280-1000.
- Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? Conformal prediction for interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.
- M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.
- Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-aware thresholding in online conformal prediction for time series. In *Communications in Computer and Information Science*, pages 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/2010/0480.
- Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.
- Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep Koopman linear embeddings. arXiv:2511.12760v2, 2025.
- Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.
- Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Salvador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026. doi: 10.1016/j.eswa.2025.129305.
- Tilman Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(2): 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.
- Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4): 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.
- Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal prediction. arXiv:2512.07770v2, 2026.
- Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for multi-horizon time series forecasting. arXiv:2604.13253v1, 2026.

190 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
191 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
192 2025. doi: 10.1609/aaai.v39i17.34029.

193 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
194 tion for multi-step forecasting. arXiv:2508.13362v3, 2025.

195 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
196 metric in conformal prediction: When shorter intervals are not better. arXiv:2601.21455v2, 2026.

197 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
198 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth
199 System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.

200 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
201 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.

202 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
203 and inventory decisions. SSRN working paper 4711817, 2024.

204 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
205 machine learning controller? arXiv:2503.14379v3, 2025.

206 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
207 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
208 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.

209 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
210 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
211 007.

212 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
213 probabilistic forecasts. arXiv:2605.28531v1, 2026.

214 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
215 arXiv:2607.26577v1, 2026.

216 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
217 ing. arXiv:2410.13115v2, 2024.

218 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
219 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
220 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.

221 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
222 inference for time series. arXiv:2502.00818v2, 2025.

223 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
224 forecasting systems. arXiv:2506.05776v3, 2026.

225 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
226 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.

227 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
228 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
229 doi: 10.1175/2009MWR2960.1.

A Supporting measurements

The primary regime. Every figure in Table 1 is measured under one deterministic adversary. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$ with $\varepsilon = 10^{-9}$ against the *deployed* threshold, and the indicator is strict. The specification is not a tie rule, and the distinction is measured rather than argued. Scoring an exact-threshold play *as a miscoverage* reproduces the primary regime on every cell of Table 1: 44 of 44 identical, all eleven arms at all four horizons. Scoring the same play under the *strict* indicator does not, and it moves every row, because the adversary’s signature move then covers rather than misses: partial adjustment at $w = 0.999$ excurses 70.80 in place of 623.70, and the failing dead band returns 0.000000 in place of 1.000000. The $\varepsilon \downarrow 0$ limit the harness specifies is the first reading. **No row is convention-free with respect to the second**, which is why the adversary is specified rather than tie-broken.

Excursions at $T = 10^4$, and the law across the arms. Partial adjustment excurses 6.30, 63.00 and 623.70 at $w = 0.9, 0.99$ and 0.999 . The last two are flat in T while Proposition 2’s bound grows like $\log T$, which is why the ratio falls with the horizon. Across the arms the excursion fits $\max_t |E_t| \approx \max\{0.5 h(T) + 0.9, 0.63/(1 - w)\}$. This is a fit to the measured arms, not a derived bound.

Realised miscoverage at the covering settings of Figure 2. At $\hat{q} \equiv +b/2$ the $\tau = 1.5$ band returns 0.100010 with $\max_t |E_t| = 11.20$, inside the bound. Under Angelopoulos et al.’s tangent integrator the same width returns 0.100001. Raising the saturator’s level to $4b$ leaves coverage intact and cuts the $w = 0.999$ excursion from 623.70 to 120.60, which is the first term of τ^* moving the edge rather than the readout changing behaviour.

B The four vocabularies

Table 2: Four literatures name the movement of a sequentially updated forecast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC, MQC/SMQC, congruence, nervousness); *jumpiness*, *convergence index*, *(in)consistency*; and *switching cost*. The row gives what each contributes.

Online conformal	Forecast stability	Verification	Switching-cost online learning
a distribution-free coverage guarantee for any bounded scorecaster, with a rate [Angelopoulos et al., 2023, Thm 1, Prop. 2], and the admissible set measured here to be tight	long-run- the scored movement functional, the readout map as post-processing [Godahewa et al., 2025, Eq. 3], and a priced decision [Genov et al., 2026, §4.2] and [Van Belle et al., 2026]	the matched-level design on real producers from 2012, with a control <i>stronger</i> than the pair: whole marginal [Pinson and Girard, 2012, Worsnop et al., 2018]	worst-case theory for the penalised problem: Thm 2, and a Thm 7 trade-off in which sublinear regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013]