
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 Conformal PID states long-run coverage for an arbitrary bounded scorecaster, and
2 a published corollary retains its rate under any predictable perturbation of the
3 deployed threshold inside an admissible radius. The opposite reading is in print too,
4 under overlapping authorship, and acted on. We make two claims. The record holds
5 both halves and we join them, conceding placement and derivation by name. And
6 at the null scorecaster the corollary’s condition is *tight*: a legal L_1 dead band on the
7 *completed* threshold leaves that set past $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, which
8 is the published radius itself there, and past it realised miscoverage is 1.000000
9 against $\alpha = 0.1$ while saturation fires every round. Leaving the set costs coverage
10 and not only the rate: partial adjustment there overruns Proposition 2’s bound
11 $42\times$ at $T = 10^6$ and still covers, 0.100004. The edge belongs to the saturator-
12 scorecaster pair, not to the readout, and elsewhere the condition is sufficient only:
13 under the equally legal $\hat{q} \equiv +b/2$ the failing band covers, and under $\hat{q} \equiv -b/2$
14 both the failing band and partial adjustment at $w = 0.999$ return 1.000000.

15 1 Introduction

16 An online conformal method emits a threshold before each observation and moves it afterwards. Fore-
17 cast stability names that movement, scores it and prices it [Tunc et al., 2013, Godahewa et al., 2025,
18 Van Belle et al., 2023, Pritularga and Kourentzes, 2024, Genov et al., 2026, Van Belle et al., 2026].
19 Varying a forecast’s temporal structure at fixed level was invented earlier, in forecast verification,
20 where the predictive marginal is matched on real producers [Gneiting et al., 2007, Pinson and Girard,
21 2012, Worsnop et al., 2018]; that coverage and mean length under-describe a deployed interval is
22 established too [Min et al., 2026, Vaze, 2026]. None of it is claimed here.

23 **First: the record carries a claim and its refutation, and this paper joins them.** Conformal PID
24 [Angelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator and asks of it
25 only that it be predictable and lie in $[-b/2, b/2]$ — “any function of the past” (Theorem 1) — so a
26 readout placed in that slot is already quantified over. Also in print, of the same method: “[s]ince
27 using a scorecaster (D-part) in C-PID breaks the theoretical coverage guarantee” [Li and Rodríguez,
28 2025] — stated and acted on in that paper’s baselines paragraph, refereed and verbatim on printed
29 page 18443. The refutation is in print too, twice and with proof: Li et al. [2025, Corollary 2] add an
30 arbitrary predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ to the deployed threshold, observe that the result
31 “has exactly the CPID error-integration form”, and conclude $|E_T| \leq ch(T) + 1 = o(T)$; the same
32 generic-slot argument carries AcMCP’s long-run coverage [Wang and Hyndman, 2024]. The first
33 and last authors of Li et al. [2025] are the authors of Li and Rodríguez [2025] and cite it: the halves
34 already share a bibliography, this is ordinary self-correction, and the joining is what this paper adds.
35 Four surfaces, queried on 20 August 2026 for the two papers’ identifiers and for “scorecaster” —
36 arXiv full text, arXiv abstracts, OpenReview, and the six works Semantic Scholar records as citing
37 Li and Rodríguez [2025] — return no document containing both the claim and the corollary. Hu et al.




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In (B) the loop closes on the integrator’s own output; in (A) on $\tilde{q}_{t+1}$, which it never computed.

[2026] still calls scorecaster choice “arbitrary” and lacking “principled guidance”, while asserting valid coverage: model selection, not validity. **We claim neither the placement nor the derivation.**

Second: that admissible set has an edge, and the condition there is tight. Corollary 2 is a *retention* result, so what it leaves open is outside that set. Section 3 exhibits a legal perturbation that leaves it — an L_1 dead band on the *completed* threshold, the family characterised here — and locates the edge at $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, the corollary’s own radius at the null scorecaster. Past it the price is coverage and not only the rate, so the published sufficient condition is necessary there. Those two claims are the paper; Section 4 sets them beside four vocabularies for the same movement.

2 Setup

The producer. A forecaster fit on data before t induces a score s_t , a threshold q_t is emitted before y_t is revealed, the deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = 1\{s_t(x_t, y_t) > q_t\}$; the target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID [Angelopoulos et al., 2023] manipulates the threshold on the score scale,

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

their iteration (5), in which r_t saturates: condition (4) requires $|r_t(x)| \geq b$ with the sign of x once $|x| \geq c h(t)$, for h nonnegative, nondecreasing and sublinear. Writing $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$, Proposition 2 gives $|E_T| \leq c h(T) + 1$; boundedness is read on realised scores.

One scalar, two readouts, two places to put it. Damping the movement of a sequentially updated scalar has two standard readouts, both prior work: a quadratic cost gives linear partial adjustment towards the target, $\tilde{q}_{t+1} = (1 - \lambda)\hat{q}_{t+1}^{\text{raw}} + \lambda\tilde{q}_t$ [Gârleanu and Pedersen, 2013], and a proportional cost a no-trade band, moving only when the target is more than τ away and then pulling back by exactly τ [Constantinides, 1986, Davis and Norman, 1990]. The first is published as post-processing [Godaheewa et al., 2025] and applied to a deployed conformal threshold [Binny and Dixit, 2025, Eq. 13]. It has two positions in (1), drawn in Figure 1: (B) in the $\hat{q}_{t+1}$ slot, where both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set; and (A) downstream of the completed q_{t+1} , outside it. Placement A is measured here. *The dead band (L1) is not a stylistic alternative to the partial-adjustment (L2) form*; it is the specific perturbation family whose failure §3 characterises. Neither is safe by construction: L2 forfeits Proposition 2’s rate while covering, L1 past τ^* loses coverage, and at the legal $\hat{q}_t \equiv -b/2$ both lose it. What is tight is the radius, not a form.

The quantity, and the names it already has. We report the deployed threshold’s *travel*, $\mathcal{T}_H = \sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new: $\sum |\Delta q|$ over a predictive distribution is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls it *jumpiness*, a *convergence index* and (*in*)*consistency* [Zsótér et al., 2009, Ehret, 2010, Pappenberger et al., 2011]; applied control, IACER [Sarhadi, 2025]; dynamic-regret online learning, a path length. Table 2’s caption names ten across four vocabularies, which is why a fresh one is used here.

Harness. $\alpha = 0.1$, $b = 2$, $c = 1$, $h(t) = \log(t + 2)$, and the clipped saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$, which meets condition (4) with equality. The scorecaster is null and legal, reducing (1) to the pure error integration Proposition 2 is stated about. Neither free choice is neutral: condition (4) is a lower bound this saturator meets exactly, and Theorem 1 permits scorecasters $b/2$ either side of the null one. The adversary is specified rather than tie-broken: it plays

Table 1: Placement A: a one-scalar readout on the *completed* threshold; one deterministic adversary, one pass to $T = 10^6$. *All eleven arms run are listed*, so no covering arm is withheld. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, “Sat.” the rounds on which (4) delivers its full $\pm b$.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5

78 $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$, $\varepsilon = 10^{-9}$, against the *deployed* $\tilde{q}_t$ — the $\varepsilon \downarrow 0$ reading of “sets the score
79 at the threshold”, firing the strict indicator. Legality is Theorem 1’s $[-b/2, b/2]$ on both sides, so
80 $b/2 + b/2 = b$ is exactly tight.

81 3 The edge of the admissible set, measured

82 **The retention result is published, and this section tests it from outside.**
83 Li et al. [2025, Corollary 2] prove that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the de-
84 ployed threshold keeps (1) in error-integration form and *retains* $|E_T| \leq c h(T) + 1$; the admissible
85 radius is theirs. A downstream partial adjustment of weight w leaves it. Its excursion at $T = 10^4$
86 is 6.30/63.00/623.70 for $w = 0.9/0.99/0.999$, and for the last two it is *flat in* T while the bound
87 grows like $\log T$, so the ratio falls with the horizon: at $w = 0.999$, 61.08 at $T = 10^4$ and 42.10 at
88 $T = 10^6$; the measured law across the arms is $\max_t |E_t| \approx \max\{0.5 h(T) + 0.9, 0.63/(1 - w)\}$,
89 and five of the eleven arms exceed the bound at $T = 10^6$. The forfeit is therefore a property of the
90 smoothing *strength* and not of Placement A as such, and it buys something: the $w = 0.999$ arm cuts
91 deployed travel from 28,311 to 202. The sign is the opposite of the intuitive one. A readout on the
92 completed q_{t+1} touches neither r_t nor the integrator’s argument, so it cannot put condition (4) at risk;
93 it delays the correction the integrator has already computed, so the accumulator excurses *further*.
94 Smoothing the output does not damp E_t , it feeds it. The 623.70 is the primary regime’s figure — the
95 adversary landing $\varepsilon = 10^{-9}$ above the deployed threshold under the harness’s strict indicator — and
96 what it turns on is that ε rather than the counting convention: an adversary playing the threshold
97 *exactly* returns 70.80 under the same indicator, because its signature move then covers.

98 **The L_1 dead band leaves the radius, and the edge is measurable.** A dead band pulls the deployed
99 threshold back by exactly τ when it releases, so “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under
100 saturation the deployed value is pinned τ below the integrator’s ceiling: its realised supremum is
101 $b - \tau$ to the printed digit at every $\tau \geq 0.9$. The adversary’s best legal play is $b/2$, so coverage is
102 lost exactly when that supremum drops below it, at $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$. Figure 2 plots
103 that edge across settings. Every dead-band configuration in the variation sweep agrees, 19 of 19:
104 three widths under each of the null scorecaster, $\hat{q} \equiv +b/2$, $\hat{q} \equiv -b/2$, a saturator at level $4b$ and
105 Angelopoulos et al.’s tangent integrator, plus four wider bands at the null scorecaster. An eleven-point
106 grid at the null scorecaster *locates* that edge between $\tau = 1.000$ and 1.001 ; the other four settings are
107 bracketed at three widths and not located. **Here** $\sup r_t = b$ **and** $\hat{q} \equiv 0$, **so** $\tau^* = b/2$ — one point
108 of the admissible class, and both terms move it. At $\hat{q} \equiv +b/2$, legal, the same $\tau = 1.5$ band covers
109 at $T = 10^6$ (0.100010, $\max_t |E_t| = 11.20$, inside the bound); at $\hat{q} \equiv -b/2$, equally legal, $\tau^* = 0$
110 and $\tau = 0.5$ returns 1.000000, as does partial adjustment at $w = 0.999$. Raising the saturator’s level
111 to $4b$, admissible since (4) is a lower bound, raises $\sup_x r_t(x)$ and τ^* with it and cuts 623.70 to
112 120.60; under the unbounded tangent integrator Angelopoulos et al. [2023] report using in all their
113 experiments, $\sup_x r_t(x)$ is unbounded and so is τ^* , the three widths run there all cover, $\tau = 1.5$ at
114 0.100001, and 623.70 falls to 20.10, both at $T = 10^6$. **The method’s own integrator therefore**
115 **sits inside the set at every width run here.** Outside it the transition is a step: $\tau = 1.000$ covers
116 at $T = 10^5$ while $\tau = 1.001$ returns 1.000000, and past τ^* the threshold sits permanently below




Figure 2: The dead-band coverage cliff, and the fact that its edge moves. **(a)** Realised miscoverage against dead-band width τ under five admissible $(r_t, \hat{q})$ pairs. Each pair’s $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$ is dashed in its own colour: 1 at the null scorecaster, 2 at $\hat{q} \equiv +b/2$, 0 at $\hat{q} \equiv -b/2$, 7 at a saturator of level $4b$, and unbounded under Angelopoulos et al.’s tangent integrator (arrow). All 19 dead-band runs land on the side of their own τ^* that the law predicts, and the vertical dashes are the law’s prediction rather than four located edges: only the null scorecaster’s edge is bracketed from both sides. **(b)** An eleven-point grid at the null scorecaster *locates* that edge between $\tau = 1.000$ and 1.001 , with no intermediate regime. Horizons: $T = 10^6$ in (a) except the four widest bands ($\tau \geq 2$), run only to $T = 2 \times 10^5$; $T = 10^5$ in (b).

the reachable score range and $|E_t| = (1 - \alpha)t$ grows linearly, 60,747 times the bound at $T = 10^6$; $\tau = 2.0, 2.5, 3.0$ and 5.0 fail too at $T = 2 \times 10^5$, so the failing set is (τ^*, ∞) , open at the left. **The endpoint is the one row in the record where the adversary’s convention decides the answer:** at $\tau = \tau^*$ the deployed supremum equals the adversary’s own cap $b/2$, so the primary regime covers while the exact-threshold play returns 1.000000 or 0.000000 according to the tie rule. Every other row is convention-free. **And the failing arms have** $\text{Sat.} = 1.0000$: condition (4) delivers its full $\pm b$ on every round, and what fails is Proposition 2’s other hypothesis, that the threshold closing the loop *is* the integrator’s output.

What that measures is that the published condition is tight. At $\mu_t = 1$, $\hat{q}_t \equiv 0$ the measured τ^* is exactly Corollary 2’s radius, and crossing it costs coverage rather than the rate. **The claim is that necessity: a tightness result about someone else’s sufficient condition**, smaller and more precise than a forfeit of the rate by post-processing. Elsewhere the condition stays sufficient only — at $\hat{q}_t \equiv +b/2$ Corollary 2 permits radius 0 while every width run there covers — which is what makes “tight” a located claim.

The constructive reading, and it is the corollary’s rather than ours. Placement B keeps $\{\hat{q}_t\}$ inside the admissible radius, so Theorem 1 carries it already; and at Placement A, condition (4) bounds $\sup_x r_t(x)$ only from below, so raising the saturator’s level widens the band at no cost to the theorem. What the measurement adds is that the band has an edge, that the edge is computable before deployment from $\sup_x r_t(x)$ and $\sup_t \hat{q}_t$, and that crossing it is not a graceful degradation.

4 Where this sits

Four vocabularies name one quantity many ways, and the object that settles its validity — the admissible set around the deployed threshold — sits in the first of them.

What is conceded, by name. The placement is prior work: Angelopoulos et al. [2023] state coverage for any admissible integrator and any scorecaster and run a smoothing-family model in that slot themselves; Dupuy et al. [2026] give the generic argument in their Appendix A and Duerst

Table 2: Four literatures name the movement of a sequentially updated forecast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC, MQC/SMQC, congruence, nervousness); *jumpiness, convergence index, (in)consistency*; and *switching cost*. The row gives what each contributes.

Online conformal	Forecast stability	Verification	Switching-cost OL
a distribution-free coverage guarantee for any bounded scorecaster, with a rate [Angelopoulos et al., 2023, Thm 1, Prop. 2], and the admissible set measured here to be tight	long-run- the scored movement functional, the readout map as post-processing [Godaheva et al., 2025, Eq. 3], and a priced decision [Genov et al., 2026, §4.2], [Van Belle et al., 2026]	the matched-level design on real producers from 2012, with a control <i>stronger</i> than the pair: the whole marginal [Pinson and Girard, 2012, Worsnop et al., 2018]	worst-case theory for the penalised problem: Thm 2, and a Thm 7 trade-off in which sublinear regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013]

et al. [2024] adopt that scorecaster. The derivation is prior work too, and that is the concession that matters most: Li et al.’s Corollary 2 proves the retention statement in Conformal PID’s own notation [Li et al., 2025], with Wang and Hyndman [2024] an independent second instance. The smoother object, the metric arithmetic and both readout maps are conceded in Section 2, and Zheng et al. [2025], Wu et al. [2025], Chen et al. [2023], Lade et al. [2026] each place a smoother inside their own validity argument. One vocabulary out, this is integrator anti-windup with a feedforward path, where free-until-the-bound-binds is the defining specification rather than a finding [Galeani et al., 2009]. Nearest in the other direction, Gao et al. [2025] run the same recursion but gate the *model* update on the conformal score; their one non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$ for a sublinear h , supported with a plot rather than derived. That accumulation is what Section 3 measures diverging.

5 Limitations

One construction, one adversary, one saturator, and $0.63/(1 - w)$ is specific to the saturator’s level, which condition (4) bounds only from below. **The inherited guarantee is weak where it applies.** Under the tangent integrator $h(t) = t/\log t$, so the certified gap $(ch(T) + 1)/T$ decays as $O(1/\log T)$ rather than $O(1/T)$: what is forfeited was not a tight certificate. **The dead-band widths are absolute, and the failure is total only against the adversary.** τ is in score units with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$; under i.i.d. scores it returns 0.249376 at $T = 10^6$, which is not a retreat but an adversary-free confirmation of the mechanism — the threshold pins at $b - \tau = 0.5$ and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly. **The harness was rebuilt against numbers already read**, which carries less weight than a rerun done without sight of them. **Placement B is conceded, not tested. One seed, and only the null scorecaster’s edge is located**; the other four settings are bracketed at three widths. **What would overturn the claim** is a legal $(r_t, \hat{q}_t)$ pair with $\tau \leq \tau^*$ that loses coverage, or $\tau > \tau^*$ that keeps it; the sweep has neither.

References

- Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages 741–763, 2013.
- Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal pid control for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. doi: 10.52202/075280-1000.
- Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? conformal prediction for interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.

- 182 M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of*
183 *Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.
- 184 Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic
185 forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical
186 Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- 187 Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-
188 aware thresholding in online conformal prediction for time series. In *Communications in Computer*
189 *and Information Science*, pages 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- 190 Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational
191 rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/
192 2010/0480.
- 193 Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-
194 windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.
- 195 Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep
196 Koopman linear embeddings. arXiv:2511.12760v2, 2025.
- 197 Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction
198 costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.
- 199 Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Sal-
200 vador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online
201 optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026.
202 doi: 10.1016/j.eswa.2025.129305.
- 203 Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration
204 and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(2):
205 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.
- 206 Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador
207 García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4):
208 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.
- 209 Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal
210 prediction. arXiv:2512.07770v2, 2026.
- 211 Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for
212 multi-horizon time series forecasting. arXiv:2604.13253v1, 2026.
- 213 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
214 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
215 2025. doi: 10.1609/aaai.v39i17.34029.
- 216 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
217 tion for multi-step forecasting. arXiv:2508.13362v3, 2025.
- 218 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
219 metric in conformal prediction: When shorter intervals are not better. arXiv:2601.21455v2, 2026.
- 220 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
221 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth*
222 *System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.
- 223 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
224 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.
- 225 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
226 and inventory decisions. SSRN working paper 4711817, 2024.
- 227 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
228 machine learning controller? arXiv:2503.14379v3, 2025.

- 229 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
230 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
231 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.
- 232 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
233 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
234 007.
- 235 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
236 probabilistic forecasts. arXiv:2605.28531v1, 2026.
- 237 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
238 arXiv:2607.26577v1, 2026.
- 239 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
240 ing. arXiv:2410.13115v2, 2024.
- 241 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
242 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
243 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.
- 244 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
245 inference for time series. arXiv:2502.00818v2, 2025.
- 246 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
247 forecasting systems. arXiv:2506.05776v3, 2026.
- 248 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
249 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.
- 250 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
251 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
252 doi: 10.1175/2009MWR2960.1.