
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 A printed claim that the scorecaster breaks online conformal coverage meets its
2 refutation by overlapping authors; this paper joins them and locates the edge
3 of the set that refutation relies on. Conformal PID states long-run coverage for
4 any bounded scorecaster; a corollary keeps it for predictable perturbations of the
5 deployed threshold inside an admissible radius. The claim was acted on; placement
6 and derivation are conceded by name. At the null scorecaster that condition is *tight*.
7 A legal L_1 dead band on the *completed* threshold leaves that set at its radius, and
8 past it miscoverage is 1.000000 against $\alpha = 0.1$. Coverage goes, not only the rate.
9 The edge belongs to the saturator–scorecaster pair, not the readout; elsewhere the
10 condition is sufficient only. Under the equally legal $\hat{q} \equiv +b/2$ the failing band
11 covers; at $\hat{q} \equiv -b/2$ both it and partial adjustment at $w = 0.999$ miscover.

1 Introduction

13 A refereed paper states that a scorecaster breaks Conformal PID’s coverage guarantee. A later one,
14 by the same first and last authors, proves in the same notation that a bounded perturbation of the
15 deployed threshold does not. The later cites the earlier, so this is ordinary self-correction, and the
16 joining is what this paper adds.

17 An online conformal method [Angelopoulos and Bates, 2021] emits a threshold before each obser-
18 vation and moves it. Forecast stability names, scores and prices that movement [Tunc et al., 2013,
19 Godahewa et al., 2025, Van Belle et al., 2023, Pritularga and Kourentzes, 2024, Genov et al., 2026,
20 Van Belle et al., 2026]. Varying temporal structure at fixed level came earlier, in forecast verification,
21 matching predictive marginals on real producers [Gneiting et al., 2007, Pinson and Girard, 2012,
22 Worsnop et al., 2018]. Coverage and mean length under-describe a deployed interval [Min et al.,
23 2026, Vaze, 2026]. None of it is claimed here.

24 **The record carries a claim and its refutation; this paper joins them.** Conformal PID [An-
25 gelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator, asking only that it be
26 predictable and lie in $[-b/2, b/2]$: “any function of the past” (Theorem 1). Such a readout is already
27 quantified over. The claim reads “[s]ince using a scorecaster (D-part) in C-PID breaks the theoretical
28 coverage guarantee” [Li and Rodríguez, 2025], stated and acted on in its baselines paragraph, ver-
29 batim on refereed printed page 18443. It is refuted twice, with proof. Li et al. [2025, Corollary 2]
30 add any predictable d_t , $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ with $\mu_t \in [0, 1]$ predictable, to the deployed threshold;
31 since $\hat{q}_t + d_t$ still lies in $[-b/2, b/2]$ it is itself a legal scorecaster, so the result “has exactly the CPID
32 error-integration form” and $|E_T| = o(T)$. The same generic-slot argument carries AcMCP’s long-run
33 coverage [Wang and Hyndman, 2024]. Both identifiers and “scorecaster” were searched across arXiv
34 full text, arXiv abstracts, OpenReview, and Semantic Scholar’s six citers of Li and Rodríguez [2025]:
35 none returns both the claim and the corollary. Hu et al. [2026] still call scorecaster choice “arbitrary”
36 and lacking “principled guidance” while asserting valid coverage: model selection, not validity. **We**
37 **claim neither the placement nor the derivation.**




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In **(B)** the loop closes on the integrator’s own output, and both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set. In **(A)** it closes on $\tilde{q}_{t+1}$, which it never computed, and the deployed sequence leaves that set.

38 **Where that refutation stops, the condition is tight.** Corollary 2 is a *retention* result, silent outside
 39 its radius, and the question it leaves is whether that radius is sharp. Section 3 exhibits a legal
 40 perturbation leaving it: the L_1 dead-band family characterised here on the *completed* threshold,
 41 whose edge is the corollary’s radius at the null scorecaster. Past it coverage goes, not only the rate: the
 42 published sufficient condition is necessary there. Section 4 sets both claims beside four vocabularies
 43 for the same movement.

44 2 Setup

45 **The producer.** A forecaster fit before t induces a score s_t , and a threshold q_t is emitted before
 46 y_t is revealed. The deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = \mathbf{1}\{s_t(x_t, y_t) > q_t\}$. The
 47 target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID
 48 [Angelopoulos et al., 2023] manipulates the threshold on the score scale, generalizing the running-
 49 error feedback that adaptive conformal inference [Gibbs and Candès, 2021] and its distribution-shift
 50 variant DtACI [Gibbs and Candès, 2024] already apply directly to the threshold:

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

51 Angelopoulos et al.’s iteration (5). Its r_t saturates: condition (4) requires $|r_t(x)| \geq b$ with the sign of
 52 x once $|x| \geq c h(t)$, for h nonnegative, nondecreasing and sublinear. With $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$,
 53 Proposition 2 gives $|E_T| \leq c h(T) + 1$, read on realised scores.

54 **One scalar, two readouts, two places.** Damping a sequentially updated scalar has two standard
 55 readouts, both prior work. A quadratic cost gives linear partial adjustment $\tilde{q}_{t+1} = (1 - \lambda)\hat{q}_{t+1}^{\text{raw}} + \lambda\hat{q}_t$
 56 [Gårleanu and Pedersen, 2013], published as post-processing [Godahewa et al., 2025] and applied
 57 to a deployed threshold [Binny and Dixit, 2025, Eq. 13]. A proportional cost gives a no-trade band,
 58 moving only when the target is more than τ away, then pulling back exactly τ [Constantinides,
 59 1986, Davis and Norman, 1990]. Either readout has two positions in (1), drawn in Figure 1; **(A)** is
 60 measured here. *The dead band (L_1) is not a stylistic alternative to the partial-adjustment (L_2) form;*
 61 *it is the family whose failure Section 3 characterises.* Neither is safe by construction: L_2 forfeits
 62 Proposition 2’s rate while covering, L_1 past the radius loses coverage, and at the legal $\hat{q} \equiv -b/2$ both
 63 lose it. What is tight is the radius, not a form.

64 **The quantity, and the names it already has.** We report the deployed threshold’s *travel*, $\mathcal{T}_H =$
 65 $\sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new. Over a predictive
 66 distribution $\sum |\Delta q|$ is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls it
 67 *jumpiness*, a *convergence index*, (in)consistency [Zsótér et al., 2009, Ehret, 2010, Pappenberger
 68 et al., 2011]; applied control, IACER [Sarhadi, 2025]; dynamic-regret online learning, a path length.
 69 Table 3’s caption names ten, so a fresh one is used here.

70 **Harness.** The harness fixes $\alpha = 0.1$ and $b = 2$. Its saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$,
 71 with $c = 1$ and $h(t) = \log(t + 2)$, meets (4) with equality. The scorecaster is null and legal, reducing
 72 (1) to the pure error integration Proposition 2 addresses. Neither choice is neutral: (4) is a lower
 73 bound met exactly here, and Theorem 1 permits scorecasters $b/2$ either side of null, predictable
 74 but not constant in t . The adversary is specified, not tie-broken. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$,
 75 $\varepsilon = 10^{-9}$, against the *deployed* $\tilde{q}_t$: the $\varepsilon \downarrow 0$ reading of “sets the score at the threshold”, firing
 76 the strict indicator. Legality is Theorem 1’s $[-b/2, b/2]$ on both sides, so $b/2 + b/2 = b$ is exactly

Table 1: Placement A: a one-scalar readout on the *completed* threshold, under one deterministic adversary, on one pass to $T = 10^6$; all eleven arms run are listed. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, and “Sat.” the fraction of rounds on which (4) delivers its full $\pm b$. Every number is the *primary* regime’s, defined in Appendix A, which also gives the $T = 10^4$ excursions and the fitted excursion law.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5




Figure 2: Realised miscoverage against dead-band width τ at the null scorecaster, $\hat{q} \equiv 0$ with a saturator of level b , under the *specified* adversary, where the upward boundary τ^{*+} of (2) is 1. **(a)** The full sweep: filled markers run to $T = 10^6$ and the four widest bands, drawn open, to $T = 2 \times 10^5$; blue covers, orange forfeits, and the dashed rule is τ^{*+} . **(b)** The same setting at $T = 10^5$ over eleven measured widths, placed by $\tau - \tau^{*+}$ on a symlog axis whose linear zone is $|\tau - \tau^{*+}| \leq 10^{-3}$, which magnifies the neighbourhood of τ^{*+} : visual spacing here is not proportional to true distance in τ . Unlabelled minor ticks mark $\tau = 0.95$ and 1.05 . The grey band spans the widest covering width to the narrowest failing one; Figure 3 gives the other four $(r_t, \hat{q})$ settings. The same eleven widths under the *mirror* adversary, which locates τ^{*-} , are in Appendix A.

77 tight. That specification names one member of a class: any adapted $s_t \in [-b/2, b/2]$ is legal, the
78 constant $s_t \equiv -b/2$ included. Every number in Table 1 is measured against the one adversary above;
79 Section 3’s two boundaries are quantified over the class, and they differ.

80 3 The edge of the admissible set, located and proved

81 **The retention result is published, and tested from outside.** Li et al. [2025, Corollary 2] prove
82 that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the deployed threshold keeps (1) in error-integration
83 form and *retains* $|E_T| \leq c h(T) + 1$: the radius is theirs. A downstream partial adjustment of weight
84 w leaves it; Table 1 prices that departure, and five of its eleven arms exceed the bound at $T = 10^6$.
85 For $w = 0.99$ and 0.999 the excursion is *flat* in T while the bound grows like $\log T$. At $w = 0.999$
86 the ratio is 61.08 at $T = 10^4$ and 42.10 at $T = 10^6$. The forfeit is a property of the smoothing
87 *strength*, not Placement A, and it buys something: travel falls from 28,311 to 202. The forfeit’s sign
88 is counter-intuitive. A readout on the completed q_{t+1} touches neither r_t nor the integrator’s argument,
89 so it cannot put (4) at risk. It delays a correction already computed, and the accumulator excurses
90 *further*. Smoothing the output does not damp E_t ; it feeds it. That 623.70 turns on the adversary’s ε ,
91 not the counting convention.

92 **The L_1 dead band leaves the radius, and the edge is a theorem.** A dead band pulls the deployed
93 threshold back by exactly τ when it releases and holds it otherwise, so $|\hat{q}_{t+1} - q_{t+1}| \leq \tau$ at every

round and “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under saturation the deployed value is pinned τ below the integrator’s ceiling, and its realised supremum is $b - \tau$ to the printed digit at every $\tau \geq 0.9$. Outside the radius two levels decide the answer and condition (4) names only one of them. Write $A_t^+ = \sup_x r_t(x)$ and $A_t^- = \inf_x r_t(x)$ for the integrator’s extremes, and $\Lambda_t^+ = \inf_{x \geq ch(t)} r_t(x)$, $\Lambda_t^- = \sup_{x \leq -ch(t)} r_t(x)$ for the levels it is *guaranteed* to reach once the accumulator clears the radius. Condition (4) says exactly $\Lambda_t^+ \geq b$ and $\Lambda_t^- \leq -b$; that $\Lambda_t^\pm = A_t^\pm$, that the integrator reaches its extremes where (4) puts them, is a further hypothesis and not one (4) supplies. Set

$$\tau^{*\pm} = \sup_t (|A_t^\pm| \pm \hat{q}_{t+1}) - b/2, \quad \sigma^\pm = \inf_t (|\Lambda_t^\pm| \pm \hat{q}_{t+1}) - b/2. \quad (2)$$

Failure, over the whole admissible class. If $\tau > \tau^{*+}$ then $\hat{q}_{t+1} + A_t^+ - \tau < b/2$ at every round, so $\hat{q}_{t+1} \leq \max(\hat{q}_t, q_{t+1} - \tau)$ makes $\{\hat{q} < b/2\}$ absorbing: a release upward lands below $b/2$, and a hold or a release downward cannot lift the deployed value at all. Against $s_t = \text{clip}(\hat{q}_t + \varepsilon, \pm b/2)$ the indicator is then $1\{\hat{q}_t < b/2\}$ and every round misses. If $\tau \geq \tau^{*-}$ the constant legal play $s_t \equiv -b/2$ makes $\{\hat{q} \geq -b/2\}$ absorbing by the mirror inequality, so no round can miss and the realised rate is 0 rather than α . Either way long-run coverage is lost, with $|E_T|$ linear in T . Nothing here asks $\hat{q}$ to be constant, asks the extremes to be attained anywhere, or asks $A_t^\pm$ to be finite beyond what the hypothesis already forces. **Retention, keyed to the reach rather than the supremum.** If $\tau \leq \sigma^+$ then $E_{t-1} \geq ch(t-1)$ forces $q_t \geq b/2 + \tau$ and so $\hat{q}_t \geq b/2$, whence $\text{err}_t = 0$ against every legal adversary; if $\tau < \sigma^-$ the mirror forces $\text{err}_t = 1$ below $-ch(t-1)$. The accumulator cannot then leave $[-ch(t), ch(t)]$ by more than one round, and $|E_T| \leq ch(T) + 1$: Proposition 2’s own bound, retained verbatim rather than weakened to $o(T)$. Using condition (4) alone this already certifies $\tau \leq b/2 + \inf_t \hat{q}_t$ and $\tau < b/2 - \sup_t \hat{q}_t$, which for a constant scorecaster is Corollary 2’s radius at $\mu_t = 1$.

The two directions meet, with no gap, on the sub-class the five settings inhabit. If $\hat{q}_t \equiv \hat{q}$, if $\Lambda_t^\pm = A_t^\pm$, and if neither depends on t , then $\sigma^\pm = \tau^{*\pm}$ and coverage is retained against every legal adapted adversary if and only if $\tau \leq \tau^{*+}$ and $\tau < \tau^{*-}$. For a symmetric saturator that is one line: coverage if and only if $\tau \leq \sup_x r_t(x) - b/2 - |\hat{q}|$, with the endpoint included only when $\hat{q} < 0$. Every setting run here is in that sub-class, so the sharp form applies to all five. **What was printed before is the failure half.** τ^{*+} is the rule Figures 2 and 3 draw, and all 19 dead-band runs agree with it read against the adversary the harness plays. It is not a retention condition, and Table 2 separates the two with legal objects only. The supremum can sit where the dynamics never go: $|E_t| \leq (1 - \alpha)t$, so a saturator equal to the harness’s inside $|x| \leq t^2$ and to $\pm 10b$ outside it satisfies (4), reads $\tau^{*+} = 19$, and fails at $\tau = 1.001$. A legal time-varying $\hat{q}$ splits $\sup_t$ from $\inf_t$: $\hat{q}_t = b/2$ on t a power of two and 0 elsewhere reads $\tau^{*+} = 2$ against $\sigma^+ = 1$. And τ^{*-} was missing entirely.

The mirror boundary is already measured, on this paper’s own grid. Under the strict indicator both $s_t \equiv -b/2$ and $s_t = \text{clip}(\hat{q}_t, \pm b/2)$ give $\text{err}_t = 1\{\hat{q}_t < -b/2\}$, so they are one error process. The τ -cliff run’s exact-threshold column under that indicator is therefore the mirror adversary, measured at the null scorecaster on eleven widths: it returns between 0.09992 and 0.09997 at $\tau = 0.5$ through 0.999, and 0.000000 with $\max_t |E_t| = 10,000 = \alpha T$ at $\tau = 1$ and at every wider band. That is $\tau^{*-} = 1$, closed on the failing side, exactly where (2) puts it. **So the failure criterion is $|E_T|$, not the miscoverage rate.** Both 1.000000 and 0.000000 are failures of the same order, $|E_T| = (1 - \alpha)T$ and αT ; and once $\hat{q}$ varies the rate stops being a tell at all. At $\tau = 1.9$ with the power-of-two scorecaster the rate is 0.110960 against $\alpha = 0.1$ while $\max_t |E_t|$ is 84,745.70 at $T = 10^6$, so a reader watching only the rate would call that coverage.

Both degenerate edges, and what they need. An unbounded supremum is not on its own enough for $\tau^{*+} = \infty$ to mean what it looks like: the same unreachable-set construction with $r_t(x) = x$ past $|x| > t^2$ has no finite supremum and fails at $\tau = 1.5$. What the tangent family has, and that construction does not, is a genuine reach: $\Lambda_t(X) = \tan(\arctan(b)X/(ch(t)))$, so the accumulator level at which the deployed threshold is forced to $b/2$ is $\kappa(\tau)ch(t)$ with $\kappa(\tau) = \arctan(b/2 + \tau - \hat{q})/\arctan(b)$, finite at every finite τ . Coverage is retained at every width, as measured; but $\kappa(\tau) > 1$ once $\tau > b/2 + \hat{q}$, so Proposition 2’s own *constant* is forfeited past $\tau = b/2$ while coverage is not. At $\tau = 1.5$, $T = 10^6$ the excursion is 15.10 against Proposition 2’s 14.8155 and against $\kappa(1.5)ch(T) + 1 = 15.8530$. At the other edge, $\sup_t \hat{q}_t = -b/2$ forces $\hat{q} \equiv -b/2$, since every $\hat{q}_t$

Table 2: Legal $(r_t, \hat{q})$ configurations on which the printed τ^* over-certifies, each under two legal adversaries. “Proved window” is where (2) certifies retention against every legal adversary; between it and the failure boundaries $\tau^{*\pm}$ neither direction is claimed, and the rows built on a supremum the dynamics cannot reach or never attain sit there. Here $\emptyset$ marks a configuration for which no width is certified, $\tau = 0$ included. $T = 10^5$ except the two power-of-two rows, at $T = 10^6$. The last row carries no dead band at all.

$(r_t, \hat{q})$ configuration	τ	printed τ^*	proved window	$s_t = \hat{q}_t + \varepsilon$		$s_t \equiv -b/2$	
				miscov.	$\max_t E_t $	miscov.	$\max_t E_t $
$\hat{q} \equiv +0.4$, level b	0.599	1.4	$\tau < 0.6$	0.100010	7.20	0.099890	11.60
$\hat{q} \equiv +0.4$, level b	0.8	1.4	$\tau < 0.6$	0.100050	8.70	0.000000	10,000
$\sup_x r_t = 10b$, off the reach	1.001	19	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t$ unbounded, off the reach	1.5	∞	$\tau < 1$	1.000000	90,000	0.000000	10,000
$A^+ = 4b, A^- = -b$	1.5	7	$\tau < 1$	0.100020	4.30	0.000000	10,000
level $4b, \hat{q} \equiv 0$	6.999	7	$\tau < 7$	0.100070	12.20	0.100050	11.60
level $4b, \hat{q} \equiv 0$	7.0	7	$\tau < 7$	0.100090	12.20	0.000000	10,000
$\sup_x r_t = 2b$, unattained	2.99	3	$\tau < 1$	0.100320	211.10	0.098010	210.50
$\sup_x r_t = 2b$, unattained	3.0	3	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t = 2b$, attained	3.0	3	$\tau < 3$	0.100000	12.20	0.000000	10,000
$\hat{q}_t = b/2$ on powers of 2	1.0	2	$\emptyset$	0.100009	14.30	0.000000	100,000
$\hat{q}_t = b/2$ on powers of 2	1.9	2	$\emptyset$	0.110960	84,745.70	0.000000	100,000
$\hat{q} \equiv +b/2$, no band	0	2	$\emptyset$	0.100000	0.90	0.000000	10,000

is at least $-b/2$ already. There $q_{t+1} - \tau \leq -b/2 + b - \tau < b/2$ for every $\tau > 0$ and every t , and $\tilde{q}_1 = 0 < b/2$, so the failing set is absorbing from round 1 with no transient: the admissible window is the single point $\tau = 0$, and $\tau = 0.5$ returns 1.000000, as does partial adjustment at $w = 0.999$.

What survives of the measurement, and what its status now is. Outside the admissible window the transition is a step. Past τ^{*+} the threshold sits permanently below the reachable score range and $|E_t| = (1 - \alpha)t$ grows linearly, 60,747 times the bound at $T = 10^6$; the wider bands $\tau = 2.0, 2.5, 3.0$ and 5.0 fail too at $T = 2 \times 10^5$ (Figure 2). **Here $\sup_x r_t(x) = b$ and $\hat{q} \equiv 0$, so $\tau^{*+} = \tau^{*-} = b/2$,** one point of the admissible class, and both terms move it. The legal $\hat{q} \equiv +b/2$ carries the $\tau = 1.5$ band inside the bound against the specified adversary, and against the mirror one no width covers there, the unsmoothed control included. Condition (4) bounds $\sup_x r_t(x)$ only from below, so raising it costs the failure direction nothing; it moves the edge too (Figure 3) and shrinks the excursion, from 623.70 to 20.10 under the tangent integrator Angelopoulos et al. [2023] report using in all their experiments, both at $T = 10^6$. **The method’s own integrator therefore sits inside the set at every width run here. The endpoint is where the two boundaries part.** The upper one is closed and the lower one is open, which the null scorecaster cannot show because τ^{*+} and τ^{*-} coincide there: four saturators with $\tau^{*+} < \tau^{*-}$ cover at $\tau = \tau^{*+}$ exactly under both adversaries, and at $\hat{q} \equiv +0.4$, where $\tau^{*-} = 0.6 < 1.4 = \tau^{*+}$, the width 0.599 covers under both while 0.6 fails under the mirror. At the baseline the two coincide at 1, so the primary regime covers there and the mirror adversary does not; Appendix A prints both readings. **The failing arms have Sat. = 1.0000.** Condition (4) delivers its full $\pm b$ on every round, and what fails is Proposition 2’s other hypothesis: the threshold closing the loop is the integrator’s output.

What that establishes is that the published condition is tight. At $\mu_t = 1, \hat{q} \equiv 0$ the boundary is exactly Corollary 2’s radius, crossing it costs coverage rather than the rate, and (2) makes that an if-and-only-if against the whole legal adversary class rather than a pattern across five settings. **The claim is that necessity: a tightness result about someone else’s sufficient condition,** smaller and more precise than a forfeit of the rate by post-processing. Elsewhere it stays sufficient only, and now provably so: under the tangent integrator Corollary 2 permits $\tau \leq b/2$ while no finite width forfeits coverage against any legal adversary, because the reach is genuine at every level. At $\hat{q} \equiv +b/2$ the looseness goes the other way and disappears: Corollary 2 permits radius 0, and (2) permits no width at all, the unsmoothed control included. That asymmetry makes “tight” a located claim.

The constructive reading, and it is the corollary’s rather than ours. Placement B keeps $\{\hat{q}_t\}$ inside the admissible radius, so Theorem 1 carries it already. What Section 3 adds is that the band has two edges, both computable before deployment from the integrator’s guaranteed reach and the scorecaster’s range, that crossing either is not a graceful degradation, and that between $\sigma^\pm$ and $\tau^{*\pm}$, which separate only when $\hat{q}$ varies in t , neither direction is settled here.

4 Where this sits

Four vocabularies name one quantity; Table 3 (App. B) names all four. **What is conceded, by name.** The placement is prior work: Angelopoulos et al. [2023] state coverage for any admissible integrator and any scorecaster, and run a smoothing-family model in that slot themselves. Dupuy et al. [2026, Appendix A] give the generic argument, and Duerst et al. [2024] adopt that scorecaster. The derivation is prior work too, and this concession matters most. Li et al. [2025, Corollary 2] prove the retention statement in Conformal PID’s own notation, with Wang and Hyndman [2024] an independent second instance. Section 2 concedes the smoother object, the metric arithmetic and both readout maps. Zheng et al. [2025], Wu et al. [2025], Chen et al. [2023], Lade et al. [2026] each place a smoother inside their own validity argument. One vocabulary out, this is integrator anti-windup with a feedforward path: free-until-the-bound-binds is its defining specification, not a finding [Galeani et al., 2009]. Nearest in the other direction, Gao et al. [2025] run the same recursion but gate the *model* update on the conformal score. Their one non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$ for a sublinear h , supported with a plot rather than derived.

5 Limitations

The evidence is one construction, one adversary, one saturator; and $0.63/(1 - w)$ is specific to the saturator’s *level*, which condition (4) bounds only from below. **The inherited guarantee is weak where it applies.** Under the tangent integrator $h(t) = t/\log t$, the certified gap $(ch(T) + 1)/T$ decays as $O(1/\log T)$ rather than $O(1/T)$. What is forfeited was not a tight certificate. **The dead-band widths are absolute, and the failure is total only against the adversary.** τ is in score units with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$. Under i.i.d. scores it returns 0.249376 at $T = 10^6$, an adversary-free confirmation of the mechanism rather than a retreat. The threshold pins at $b - \tau = 0.5$, and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly. **The harness was rebuilt against numbers already read**, which carries less weight than a rerun done without sight of them. **Placement B is conceded, not tested. The i.i.d. check above used one seed** (20260820, $T = 10^6$); **the primary, adversarial harness uses none**, and **only its null-scorecaster edge is bracketed**; the others are one-sided. **What would overturn the claim** is a legal $(r_t, \hat{q})$ pair with $\tau \leq \tau^*$ that loses coverage, or $\tau > \tau^*$ that keeps it; the sweep has neither.

References

- Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages 741–763, 2013.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. arXiv:2107.07511v6, 2021.
- Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. doi: 10.52202/075280-1000.
- Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? Conformal prediction for interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.
- M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.

- 229 Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic
230 forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical
231 Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- 232 Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-
233 aware thresholding in online conformal prediction for time series. In *Communications in Computer
234 and Information Science*, pages 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- 235 Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational
236 rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/
237 2010/0480.
- 238 Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-
239 windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.
- 240 Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep
241 Koopman linear embeddings. arXiv:2511.12760v2, 2025.
- 242 Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction
243 costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.
- 244 Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Sal-
245 vador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online
246 optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026.
247 doi: 10.1016/j.eswa.2025.129305.
- 248 Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In
249 *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021.
- 250 Isaac Gibbs and Emmanuel Candès. Conformal inference for online prediction with arbitrary
251 distribution shifts. *Journal of Machine Learning Research*, 2024.
- 252 Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration
253 and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(2):
254 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.
- 255 Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador
256 García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4):
257 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.
- 258 Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal
259 prediction. arXiv:2512.07770v2, 2026.
- 260 Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for
261 multi-horizon time series forecasting. arXiv:2604.13253v1, 2026.
- 262 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
263 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
264 2025. doi: 10.1609/aaai.v39i17.34029.
- 265 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
266 tion for multi-step forecasting. arXiv:2508.13362v3, 2025.
- 267 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
268 metric in conformal prediction: When shorter intervals are not better. arXiv:2601.21455v2, 2026.
- 269 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
270 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth
271 System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.
- 272 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
273 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.
- 274 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
275 and inventory decisions. SSRN working paper 4711817, 2024.

- 276 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
277 machine learning controller? arXiv:2503.14379v3, 2025.
- 278 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
279 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
280 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.
- 281 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
282 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
283 007.
- 284 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
285 probabilistic forecasts. arXiv:2605.28531v1, 2026.
- 286 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
287 arXiv:2607.26577v1, 2026.
- 288 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
289 ing. arXiv:2410.13115v2, 2024.
- 290 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
291 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
292 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.
- 293 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
294 inference for time series. arXiv:2502.00818v2, 2025.
- 295 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
296 forecasting systems. arXiv:2506.05776v3, 2026.
- 297 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
298 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.
- 299 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
300 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
301 doi: 10.1175/2009MWR2960.1.

302 Code and data

303 The simulator, its frozen configuration, the results/ files behind every figure and table, the gen-
304 erators that read them, and this project’s audit trail across sessions are at [https://\[REPOSITORY](https://[REPOSITORY LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION])
305 [LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION\]](https://[REPOSITORY LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION]).

306 A Supporting measurements

307 **The primary regime.** Every number in Table 1 is measured under one deterministic adversary,
308 which consumes no randomness and needs no seed. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$ with $\varepsilon = 10^{-9}$
309 against the *deployed* threshold, and the indicator is strict. The specification is not a tie rule, and
310 the distinction is measured rather than argued. Scoring an exact-threshold play as a *miscoverage*
311 reproduces every quantity Table 1 reports: 44 of 44 arm-horizon cells identical, all eleven arms at
312 all four horizons. Scoring the same play under the *strict* indicator does not, and it moves every
313 row, because the adversary’s signature move then covers rather than misses. Partial adjustment at
314 $w = 0.999$ excurses 70.80 in place of 623.70, and the failing dead band returns 0.000000 in place of
315 1.000000. The $\varepsilon \downarrow 0$ limit the harness specifies is the first reading. **No row is convention-free with**
316 **respect to the second**, which is why the adversary is specified rather than tie-broken. At $\tau = \tau^{*+}$
317 itself the two readings part: the strict indicator returns 0.000000 where scoring the exact play as a
318 miscoverage returns 1.000000. **Neither of those is coverage.** They are $|E_T| = \alpha T$ and $(1 - \alpha)T$,
319 failures of the same order and opposite sign. Under the strict indicator the exact-threshold play scores
320 $\text{err}_t = \mathbf{1}\{\tilde{q}_t < -b/2\}$, which is also the indicator of the constant legal play $s_t \equiv -b/2$, so that
321 column is Section 3’s *mirror adversary* rather than a second convention. It returns between 0.09992

and 0.09997 at $\tau = 0.5, 0.9, 0.95, 0.99$ and 0.999 , and 0.000000 with $\max_t |E_t| = 10,000$ at $\tau = 1, 1.001, 1.01, 1.05, 1.1$ and 1.5 . That locates $\tau^{*-} = 1$ on the same eleven widths Figure 2 uses, closed on the failing side. The convention decides which legal adversary realises the endpoint failure, not whether one occurs.

Excursions at $T = 10^4$, and where the excursion law holds. Partial adjustment excurses 6.30, 63.00 and 623.70 at $w = 0.9, 0.99$ and 0.999 . The last two are flat in T while Proposition 2's bound grows like $\log T$, which is why the ratio falls with the horizon. On the control and the four partial-adjustment arms the excursion fits $\max_t |E_t| \approx \max\{0.5 h(T) + 0.9, 0.63/(1-w)\}$ at every horizon, most loosely at $w = 0.9$, the one arm whose two branches cross inside the horizon range. This is a fit to those five arms, not a derived bound. It does not extend to the other six: the growing time constants, the running mean and the dead bands all excurse further than it predicts.




Figure 3: All 19 dead-band runs, at the five admissible $(r_t, \hat{q})$ settings: rows are settings ordered by τ^* , and the shared horizontal axis is τ . Blue markers cover at $\approx \alpha$ and orange markers return 1.000000; filled markers run to $T = 10^6$ and open markers to $T = 2 \times 10^5$, all under the *specified* adversary. The rule on each row is that setting's upward boundary $\tau^{*+} = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, printed at the right, and the tinted band is $\tau \leq \tau^{*+}$. Two rules fall off the axis and are drawn as arrows: $\tau^{*+} = 0$ at $\hat{q} \equiv -b/2$, and τ^{*+} unbounded under Angelopoulos et al.'s tangent integrator. The mirror boundary τ^{*-} of (2) is not drawn here and differs from τ^{*+} at the two non-null scorecasters.

Realised miscoverage at the covering settings of Figure 3. At $\hat{q} \equiv +b/2$ the $\tau = 1.5$ band returns 0.100010 with $\max_t |E_t| = 11.20$, inside the bound. Under Angelopoulos et al.'s tangent integrator the same width returns 0.100001. Raising the saturator's level to $4b$ leaves coverage intact and cuts the $w = 0.999$ excursion from 623.70 to 120.60. That is the first term of τ^* moving the edge, not the readout changing behaviour.

Table 3: Four literatures name the movement of a sequentially updated forecast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC, MQC/SMQC, congruence, nervousness); *jumpiness*, *convergence index*, *(in)consistency*; and *switching cost*. The row gives what each contributes.

Online conformal	Forecast stability	Verification	Switching-cost online learning
a distribution-free coverage guarantee for any bounded scorecaster, with a rate [Angelopoulos et al., 2023, Thm 1, Prop. 2], and the admissible set measured here to be tight	long-run- the scored movement functional, the readout map as post-processing [Godaheva et al., 2025, Eq. 3], and a priced decision [Genov et al., 2026, §4.2] and [Van Belle et al., 2026]	the matched-level design on real producers from 2012, with a control <i>stronger</i> than the pair: the whole marginal [Pinson and Girard, 2012, Worsnop et al., 2018]	worst-case theory for the penalised problem: Thm 2, and a Thm 7 trade-off in which sublinear regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013]