
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 A printed claim that the scorecaster breaks online conformal coverage meets its
2 refutation by overlapping authors; this paper joins them and locates the edge
3 of the set that refutation relies on. Conformal PID states long-run coverage for
4 any bounded scorecaster; a corollary keeps it for predictable perturbations of the
5 deployed threshold inside an admissible radius. The claim was acted on; placement
6 and derivation are conceded by name. At the null scorecaster that condition is *tight*.
7 A legal L_1 dead band on the *completed* threshold leaves that set at its radius, and
8 past it miscoverage is 1.000000 against $\alpha = 0.1$. Coverage goes, not only the rate.
9 The edge belongs to the saturator–scorecaster pair, not the readout; elsewhere the
10 condition is sufficient only. Under the equally legal $\hat{q} \equiv +b/2$ the failing band
11 covers; at $\hat{q} \equiv -b/2$ both it and partial adjustment at $w = 0.999$ miscover.

1 Introduction

13 A refereed paper states that a scorecaster breaks Conformal PID’s coverage guarantee. A later one,
14 by the same first and last authors, proves in the same notation that a bounded perturbation of the
15 deployed threshold does not. The later cites the earlier, so this is ordinary self-correction, and the
16 joining is what this paper adds.

17 An online conformal method [Angelopoulos and Bates, 2021] emits a threshold before each obser-
18 vation and moves it. Forecast stability names, scores and prices that movement [Tunc et al., 2013,
19 Godahewa et al., 2025, Van Belle et al., 2023, Pritularga and Kourentzes, 2024, Genov et al., 2026,
20 Van Belle et al., 2026]. Varying temporal structure at fixed level came earlier, in forecast verification,
21 matching predictive marginals on real producers [Gneiting et al., 2007, Pinson and Girard, 2012,
22 Worsnop et al., 2018]. Coverage and mean length under-describe a deployed interval [Min et al.,
23 2026, Vaze, 2026]. None of it is claimed here.

24 **The record carries a claim and its refutation; this paper joins them.** Conformal PID [An-
25 gelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator, asking only that it be
26 predictable and lie in $[-b/2, b/2]$: “any function of the past” (Theorem 1). Such a readout is already
27 quantified over. The claim reads “[s]ince using a scorecaster (D-part) in C-PID breaks the theoretical
28 coverage guarantee” [Li and Rodríguez, 2025], stated and acted on in its baselines paragraph, ver-
29 batim on refereed printed page 18443. It is refuted twice, with proof. Li et al. [2025, Corollary 2]
30 add any predictable d_t , $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ with $\mu_t \in [0, 1]$ predictable, to the deployed threshold;
31 since $\hat{q}_t + d_t$ still lies in $[-b/2, b/2]$ it is itself a legal scorecaster, so the result “has exactly the CPID
32 error-integration form” and $|E_T| = o(T)$. The same generic-slot argument carries AcMCP’s long-run
33 coverage [Wang and Hyndman, 2024]. Both identifiers and “scorecaster” were searched across arXiv
34 full text, arXiv abstracts, OpenReview, and Semantic Scholar’s six citers of Li and Rodríguez [2025]:
35 none returns both the claim and the corollary. Hu et al. [2026] still call scorecaster choice “arbitrary”
36 and lacking “principled guidance” while asserting valid coverage: model selection, not validity. **We**
37 **claim neither the placement nor the derivation.**




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In (B) the loop closes on the integrator’s own output, and both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set. In (A) it closes on $\tilde{q}_{t+1}$, which it never computed, and the deployed sequence leaves that set.

38 **Where that refutation stops, the condition is tight.** Corollary 2 is a *retention* result, silent outside
 39 its radius, and the question it leaves is whether that radius is sharp. Section 3 exhibits a legal
 40 perturbation leaving it: the L_1 dead-band family characterised here on the *completed* threshold,
 41 whose edge is the corollary’s radius at the null scorecaster. Past it coverage goes, not only the rate: the
 42 published sufficient condition is necessary there. Section 4 sets both claims beside four vocabularies
 43 for the same movement.

44 2 Setup

45 **The producer.** A forecaster fit before t induces a score s_t , and a threshold q_t is emitted before
 46 y_t is revealed. The deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = \mathbf{1}\{s_t(x_t, y_t) > q_t\}$. The
 47 target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID
 48 [Angelopoulos et al., 2023] manipulates the threshold on the score scale, where adaptive conformal
 49 inference [Gibbs and Candès, 2021] and DtACI [Gibbs and Candès, 2024] move the level,

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

50 their iteration (5). Its r_t saturates: condition (4) requires $|r_t(x)| \geq b$ with the sign of x once $|x| \geq$
 51 $c h(t)$, for h nonnegative, nondecreasing and sublinear. With $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$, Proposition 2
 52 gives $|E_T| \leq c h(T) + 1$, read on realised scores.

53 **One scalar, two readouts, two places.** Damping a sequentially updated scalar has two standard
 54 readouts, both prior work. A quadratic cost gives linear partial adjustment $\tilde{q}_{t+1} = (1 - \lambda)\hat{q}_{t+1}^{\text{raw}} + \lambda\tilde{q}_t$
 55 [Gârleanu and Pedersen, 2013], published as post-processing [Godahewa et al., 2025] and applied
 56 to a deployed threshold [Binny and Dixit, 2025, Eq. 13]. A proportional cost gives a no-trade band,
 57 moving only when the target is more than τ away, then pulling back exactly τ [Constantinides,
 58 1986, Davis and Norman, 1990]. Either readout has two positions in (1), drawn in Figure 1; (A) is
 59 measured here. *The dead band (L_1) is not a stylistic alternative to the partial-adjustment (L_2) form;*
 60 *it is the family whose failure Section 3 characterises. Neither is safe by construction: L_2 forfeits*
 61 *Proposition 2’s rate while covering, L_1 past the radius loses coverage, and at the legal $\hat{q} \equiv -b/2$ both*
 62 *lose it. What is tight is the radius, not a form.*

63 **The quantity, and the names it already has.** We report the deployed threshold’s *travel*, $\mathcal{T}_H =$
 64 $\sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new. Over a predictive
 65 distribution $\sum |\Delta q|$ is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls it
 66 *jumpiness*, a *convergence index*, (in)consistency [Zsótér et al., 2009, Ehret, 2010, Pappenberger
 67 et al., 2011]; applied control, IACER [Sarhadi, 2025]; dynamic-regret online learning, a path length.
 68 Table 2’s caption names ten, so a fresh one is used here.

69 **Harness.** The harness fixes $\alpha = 0.1$ and $b = 2$. Its saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$,
 70 with $c = 1$ and $h(t) = \log(t + 2)$, meets (4) with equality. The scorecaster is null and legal, reducing
 71 (1) to the pure error integration Proposition 2 addresses. Neither choice is neutral: (4) is a lower
 72 bound met exactly here, and Theorem 1 permits scorecasters $b/2$ either side of null. The adversary is
 73 specified, not tie-broken. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$, $\varepsilon = 10^{-9}$, against the *deployed* $\tilde{q}_t$: the
 74 $\varepsilon \downarrow 0$ reading of “sets the score at the threshold”, firing the strict indicator. Legality is Theorem 1’s
 75 $[-b/2, b/2]$ on both sides, so $b/2 + b/2 = b$ is exactly tight.

Table 1: Placement A: a one-scalar readout on the *completed* threshold, under one deterministic adversary, on one pass to $T = 10^6$; all eleven arms run are listed. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, and “Sat.” the fraction of rounds on which (4) delivers its full $\pm b$. Every number is the *primary* regime’s, defined in Appendix A, which also gives the $T = 10^4$ excursions and the fitted excursion law.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5




Figure 2: Realised miscoverage against dead-band width τ at the null scorecaster, $\hat{q} \equiv 0$ with a saturator of level b , where $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$ is 1. (a) The full sweep: filled markers run to $T = 10^6$ and the four widest bands, drawn open, to $T = 2 \times 10^5$; blue covers, orange forfeits, and the dashed rule is τ^* . (b) The same setting at $T = 10^5$ over eleven measured widths, placed by $\tau - \tau^*$ on a symlog axis whose linear zone is $|\tau - \tau^*| \leq 10^{-3}$, which magnifies the neighbourhood of τ^* : visual spacing here is not proportional to true distance in τ . Unlabelled minor ticks mark $\tau = 0.95$ and 1.05 . The grey band spans the widest covering width to the narrowest failing one; Figure 3 gives the other four $(r_t, \hat{q})$ settings.

3 The edge of the admissible set, measured

The retention result is published, and tested from outside. Li et al. [2025, Corollary 2] prove that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the deployed threshold keeps (1) in error-integration form and retains $|E_T| \leq c h(T) + 1$: the radius is theirs. A downstream partial adjustment of weight w leaves it; Table 1 prices that departure, and five of its eleven arms exceed the bound at $T = 10^6$. For $w = 0.99$ and 0.999 the excursion is *flat* in T while the bound grows like $\log T$. At $w = 0.999$ the ratio is 61.08 at $T = 10^4$ and 42.10 at $T = 10^6$. The forfeit is a property of the smoothing *strength*, not Placement A, and it buys something: travel falls from 28,311 to 202. The forfeit’s sign is counter-intuitive. A readout on the completed q_{t+1} touches neither r_t nor the integrator’s argument, so it cannot put (4) at risk. It delays a correction already computed, and the accumulator excurses *further*. Smoothing the output does not damp E_t ; it feeds it. That 623.70 turns on the adversary’s ε , not the counting convention.

The L_1 dead band leaves the radius, and the edge is measurable. A dead band pulls the deployed threshold back by exactly τ when it releases, so “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under saturation the deployed value is pinned τ below the integrator’s ceiling, and its realised supremum is $b - \tau$ to the printed digit at every $\tau \geq 0.9$. The adversary’s best legal play is $b/2$, so coverage is lost when that supremum drops below it, at $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$. All 19 dead-band runs agree with that law (Figure 3); only the null scorecaster’s edge is located (Figure 2), the others one-sided. **Here** $\sup_x r_t(x) = b$ and $\hat{q} \equiv 0$, so $\tau^* = b/2$, one point of the admissible class, and both terms move it. The legal $\hat{q} \equiv +b/2$ carries the $\tau = 1.5$ band inside the bound. The equally

96 legal $\hat{q} \equiv -b/2$ drops the edge to 0, and there $\tau = 0.5$ returns 1.000000, as does partial adjustment
97 at $w = 0.999$. Condition (4) bounds $\sup_x r_t(x)$ only from below, so raising it costs the theorem
98 nothing; it moves the edge too (Figure 3) and shrinks the excursion. Under the unbounded tangent
99 integrator Angelopoulos et al. [2023] report using in all their experiments, τ^* is unbounded and all
100 three widths there cover. The excursion is 20.10 rather than 623.70, both at $T = 10^6$. **The method’s**
101 **own integrator therefore sits inside the set at every width run here.** Outside the set the transition
102 is a step. Past τ^* the threshold sits permanently below the reachable score range and $|E_t| = (1 - \alpha)t$
103 grows linearly, 60,747 times the bound at $T = 10^6$. The wider bands $\tau = 2.0, 2.5, 3.0$ and 5.0
104 fail too at $T = 2 \times 10^5$ (Figure 2), so the failing set is (τ^*, ∞) , open at the left. **The endpoint is**
105 **where the adversary’s specification decides the answer.** At $\tau = \tau^*$ the deployed supremum equals
106 the adversary’s own cap $b/2$, so the clip forces an exact play and the reading settles it: the primary
107 regime covers. Appendix A separates the two readings and prints both. **The failing arms have**
108 **Sat. = 1.0000.** Condition (4) delivers its full $\pm b$ on every round, and what fails is Proposition 2’s
109 other hypothesis: the threshold closing the loop *is* the integrator’s output.

110 **What that measures is that the published condition is tight.** At $\mu_t = 1$, $\hat{q} \equiv 0$ the measured
111 τ^* is exactly Corollary 2’s radius, and crossing it costs coverage rather than the rate. **The claim is**
112 **that necessity: a tightness result about someone else’s sufficient condition,** smaller and more
113 precise than a forfeit of the rate by post-processing. Elsewhere it stays sufficient only: at $\hat{q} \equiv +b/2$
114 Corollary 2 permits radius 0 while every width run there covers. That asymmetry makes “tight” a
115 located claim.

116 **The constructive reading, and it is the corollary’s rather than ours.** Placement B keeps $\{\hat{q}_t\}$
117 inside the admissible radius, so Theorem 1 carries it already. What the measurement adds is that the
118 band has an edge, computable before deployment from $\sup_x r_t(x)$ and $\sup_t \hat{q}_t$, and that crossing it is
119 not a graceful degradation.

120 4 Where this sits

121 Four vocabularies name one quantity; Table 2 (App. B) names all four. **What is conceded, by**
122 **name.** The placement is prior work: Angelopoulos et al. [2023] state coverage for any admissible
123 integrator and any scorecaster, and run a smoothing-family model in that slot themselves. Dupuy
124 et al. [2026, Appendix A] give the generic argument, and Duerst et al. [2024] adopt that scorecaster.
125 The derivation is prior work too, and this concession matters most. Li et al. [2025, Corollary 2]
126 prove the retention statement in Conformal PID’s own notation, with Wang and Hyndman [2024] an
127 independent second instance. Section 2 concedes the smoother object, the metric arithmetic and both
128 readout maps. Zheng et al. [2025], Wu et al. [2025], Chen et al. [2023], Lade et al. [2026] each place a
129 smoother inside their own validity argument. One vocabulary out, this is integrator anti-windup with a
130 feedforward path: free-until-the-bound-binds is its defining specification, not a finding [Galeani et al.,
131 2009]. Nearest in the other direction, Gao et al. [2025] run the same recursion but gate the *model*
132 update on the conformal score. Their one non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$
133 for a sublinear h , supported with a plot rather than derived.

134 5 Limitations

135 **The evidence is one construction, one adversary, one saturator;** and $0.63/(1 - w)$ is specific to
136 the saturator’s *level*, which condition (4) bounds only from below. **The inherited guarantee is weak**
137 **where it applies.** Under the tangent integrator $h(t) = t/\log t$, the certified gap $(ch(T) + 1)/T$
138 decays as $O(1/\log T)$ rather than $O(1/T)$. What is forfeited was not a tight certificate. **The dead-**
139 **band widths are absolute, and the failure is total only against the adversary.** τ is in score units
140 with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$. Under i.i.d. scores it returns 0.249376 at $T = 10^6$, an
141 adversary-free confirmation of the mechanism rather than a retreat. The threshold pins at $b - \tau = 0.5$,
142 and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly. **The harness was rebuilt**
143 **against numbers already read,** which carries less weight than a rerun done without sight of them.
144 **Placement B is conceded, not tested. The i.i.d. check above used one seed,** 20260820 at $T = 10^6$;
145 the primary, adversarial harness behind Table 1 and the figures uses none. **Only the null scorecaster’s**
146 **edge is bracketed** by that (seed-free) dead-band sweep; the others are one-sided. **What would**

147 **overturn the claim** is a legal $(r_t, \hat{q})$ pair with $\tau \leq \tau^*$ that loses coverage, or $\tau > \tau^*$ that keeps it;
 148 the sweep has neither.

149 References

- 150 Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan
 151 Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness
 152 and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages
 153 741–763, 2013.
- 154 Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and
 155 distribution-free uncertainty quantification. arXiv:2107.07511v6, 2021.
- 156 Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control
 157 for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS
 158 2023)*, 2023. doi: 10.52202/075280-1000.
- 159 Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? Conformal prediction for
 160 interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- 161 Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng
 162 Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference
 163 for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge
 164 Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- 165 George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political
 166 Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.
- 167 M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of
 168 Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.
- 169 Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic
 170 forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical
 171 Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- 172 Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-
 173 aware thresholding in online conformal prediction for time series. In *Communications in Computer
 174 and Information Science*, pages 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- 175 Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational
 176 rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/
 177 2010/0480.
- 178 Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-
 179 windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.
- 180 Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep
 181 Koopman linear embeddings. arXiv:2511.12760v2, 2025.
- 182 Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction
 183 costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.
- 184 Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Sal-
 185 vador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online
 186 optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026.
 187 doi: 10.1016/j.eswa.2025.129305.
- 188 Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In
 189 *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021.
- 190 Isaac Gibbs and Emmanuel Candès. Conformal inference for online prediction with arbitrary
 191 distribution shifts. *Journal of Machine Learning Research*, 2024.

192 Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration
193 and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(2):
194 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.

195 Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador
196 García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4):
197 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.

198 Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal
199 prediction. arXiv:2512.07770v2, 2026.

200 Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for
201 multi-horizon time series forecasting. arXiv:2604.13253v1, 2026.

202 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
203 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
204 2025. doi: 10.1609/aaai.v39i17.34029.

205 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
206 tion for multi-step forecasting. arXiv:2508.13362v3, 2025.

207 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
208 metric in conformal prediction: When shorter intervals are not better. arXiv:2601.21455v2, 2026.

209 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
210 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth
211 System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.

212 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
213 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.

214 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
215 and inventory decisions. SSRN working paper 4711817, 2024.

216 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
217 machine learning controller? arXiv:2503.14379v3, 2025.

218 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
219 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
220 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.

221 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
222 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
223 007.

224 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
225 probabilistic forecasts. arXiv:2605.28531v1, 2026.

226 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
227 arXiv:2607.26577v1, 2026.

228 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
229 ing. arXiv:2410.13115v2, 2024.

230 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
231 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
232 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.

233 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
234 inference for time series. arXiv:2502.00818v2, 2025.

235 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
236 forecasting systems. arXiv:2506.05776v3, 2026.

237 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
 238 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.

239 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
 240 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
 241 doi: 10.1175/2009MWR2960.1.

242 Code and data

243 The simulator, its frozen configuration, the results/ files behind every figure and table, the gen-
 244 erators that read them, and this project’s audit trail across sessions are at <https://github.com/pa1aash/Turnover-Blind>.
 245

246 A Supporting measurements

247 **The primary regime.** Every number in Table 1 is measured under one deterministic adversary,
 248 which consumes no randomness and needs no seed. It plays $s_t = \text{clip}(\hat{q}_t + \varepsilon, \pm b/2)$ with $\varepsilon = 10^{-9}$
 249 against the *deployed* threshold, and the indicator is strict. The specification is not a tie rule, and
 250 the distinction is measured rather than argued. Scoring an exact-threshold play *as a miscoverage*
 251 reproduces every quantity Table 1 reports: 44 of 44 arm-horizon cells identical, all eleven arms at
 252 all four horizons. Scoring the same play under the *strict* indicator does not, and it moves every
 253 row, because the adversary’s signature move then covers rather than misses. Partial adjustment at
 254 $w = 0.999$ excurses 70.80 in place of 623.70, and the failing dead band returns 0.000000 in place of
 255 1.000000. The $\varepsilon \downarrow 0$ limit the harness specifies is the first reading. **No row is convention-free with**
 256 **respect to the second**, which is why the adversary is specified rather than tie-broken. At $\tau = \tau^*$
 257 itself the two readings part: the strict indicator returns 0.000000 where scoring the exact play as a
 258 miscoverage returns 1.000000.

259 **Excursions at $T = 10^4$, and where the excursion law holds.** Partial adjustment excurses 6.30,
 260 63.00 and 623.70 at $w = 0.9, 0.99$ and 0.999 . The last two are flat in T while Proposition 2’s
 261 bound grows like $\log T$, which is why the ratio falls with the horizon. On the control and the four
 262 partial-adjustment arms the excursion fits $\max_t |E_t| \approx \max\{0.5 h(T) + 0.9, 0.63/(1 - w)\}$ at every
 263 horizon, most loosely at $w = 0.9$, the one arm whose two branches cross inside the horizon range.
 264 This is a fit to those five arms, not a derived bound. It does not extend to the other six: the growing
 265 time constants, the running mean and the dead bands all excurse further than it predicts.




Figure 3: All 19 dead-band runs, at the five admissible $(r_t, \hat{q})$ settings: rows are settings ordered by τ^* , and the shared horizontal axis is τ . Blue markers cover at $\approx \alpha$ and orange markers return 1.000000; filled markers run to $T = 10^6$ and open markers to $T = 2 \times 10^5$. The rule on each row is that setting’s $\tau^* = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, printed at the right, and the tinted band is $\tau \leq \tau^*$. Two rules fall off the axis and are drawn as arrows: $\tau^* = 0$ at $\hat{q} \equiv -b/2$, and τ^* unbounded under Angelopoulos et al.’s tangent integrator.

266 **Realised miscoverage at the covering settings of Figure 3.** At $\hat{q} \equiv +b/2$ the $\tau = 1.5$ band returns
 267 0.100010 with $\max_t |E_t| = 11.20$, inside the bound. Under Angelopoulos et al.’s tangent integrator

268 the same width returns 0.100001. Raising the saturator's level to $4b$ leaves coverage intact and cuts
269 the $w = 0.999$ excursion from 623.70 to 120.60. That is the first term of τ^* moving the edge, not the
270 readout changing behaviour.

B The four vocabularies

Table 2: Four literatures name the movement of a sequentially updated forecast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC, MQC/SMQC, congruence, nervousness); *jumpiness*, *convergence index*, *(in)consistency*; and *switching cost*. The row gives what each contributes.

Online conformal	Forecast stability	Verification	Switching-cost online learning
a distribution-free coverage guarantee for any bounded scorecaster, with a rate [Angelopoulos et al., 2023, Thm 1, Prop. 2], and the admissible set measured here to be tight	long-run- the scored movement functional, the readout map as post-processing [Godaheva et al., 2025, Eq. 3], and a priced decision [Genov et al., 2026, §4.2] and [Van Belle et al., 2026]	the matched-level design on real producers from 2012, with a control <i>stronger</i> than the pair: the whole marginal [Pinson and Girard, 2012, Worsnop et al., 2018]	worst-case theory for the penalised problem: Thm 2, and a Thm 7 trade-off in which sublinear regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013]