
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 A printed claim that the scorecaster breaks online conformal coverage meets its
2 refutation by overlapping authors; this paper joins them and locates the edge
3 of the set that refutation relies on. Conformal PID states long-run coverage for
4 any bounded scorecaster; a corollary keeps it for predictable perturbations of the
5 deployed threshold inside an admissible radius. The claim was acted on; the
6 readout’s placement and the corollary’s derivation are conceded by name. At
7 the null scorecaster that condition is *tight*. A legal dead band on the *completed*
8 threshold leaves that set at its radius, and past it miscoverage is 1.000000 against
9 $\alpha = 0.1$: coverage fails outright, not merely its rate. The edge belongs to the
10 saturator–scorecaster pair, not the readout; elsewhere the condition is sufficient
11 only. Under the equally legal $\hat{q} \equiv +b/2$ the failing band covers; at $\hat{q} \equiv -b/2$ both
12 it and partial adjustment at $w = 0.999$ miscover.

13 1 Introduction

14 A refereed paper states that a scorecaster breaks Conformal PID’s coverage guarantee. A later one,
15 by the same first and last authors, proves in the same notation that a bounded perturbation of the
16 deployed threshold does not. The later cites the earlier, so this is ordinary self-correction, and the
17 joining is what this paper adds.

18 An online conformal method [Angelopoulos and Bates, 2023] moves a threshold each round, and
19 that movement is not an unnamed quantity. Forecast stability names, scores and prices it [Tunc et al.,
20 2013, Godahewa et al., 2025, Van Belle et al., 2023, Pritularga and Kourentzes, 2024, Genov et al.,
21 2026, Van Belle et al., 2026]. Holding the level fixed and comparing the temporal structure came
22 earlier still, in forecast verification, which matches predictive marginals on real producers [Gneiting
23 et al., 2007, Pinson and Girard, 2012, Worsnop et al., 2018]. A third strand reports that the two
24 numbers usually printed of a deployed interval, coverage and mean length, under-describe it [Min
25 et al., 2026, Vaze, 2026]. None of it is claimed here.

26 **The record carries a claim and its refutation; this paper joins them.** Conformal PID [An-
27 gelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator, asking only that
28 it be predictable and lie in $[-b/2, b/2]$: “any function of the past” (Theorem 1). Such a readout
29 is already quantified over. The claim reads “[s]ince using a scorecaster (D-part) in C-PID breaks
30 the theoretical coverage guarantee” [Li and Rodríguez, 2025], stated and acted on in its baselines
31 paragraph, verbatim on refereed printed page 18443. It is refuted twice, with proof. Li et al. [2025,
32 Corollary 2] add any predictable $d_t, |d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ with $\mu_t \in [0, 1]$ predictable, to the deployed
33 threshold; since $\hat{q}_t + d_t$ still lies in $[-b/2, b/2]$ it is itself a legal scorecaster, so the result “has exactly
34 the CPID error-integration form” and Proposition 2’s own finite-horizon bound on $|E_T|$ is retained,
35 whence $|E_T| = o(T)$. The same generic-slot argument carries AcMCP’s long-run coverage [Wang
36 and Hyndman, 2024]. Four surfaces were searched for both identifiers and “scorecaster”. They are
37 arXiv full text, arXiv abstracts, OpenReview, and the six works Semantic Scholar records citing
38 Li and Rodríguez [2025]. None returns both the claim and the corollary. Hu et al. [2026] still call




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In (B) the loop closes on the integrator’s own output, and both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set. In (A) it closes on $\tilde{q}_{t+1}$, which it never computed, and nothing keeps the deployed sequence inside that set.

scorecaster choice “arbitrary” and lacking “principled guidance” while asserting valid coverage: model selection, not validity. **We claim neither the placement nor the derivation.**

Where Corollary 2 stops, its radius condition is tight. It is a *retention* result, silent outside that radius: perturbations inside the radius keep the guarantee, and perturbations outside it are not addressed either way. A sufficient condition can be conservative, so the question the corollary leaves open is whether its radius is sharp, that is, whether something legal just outside it actually fails. Section 3 exhibits a legal perturbation leaving it: the L_1 dead-band family characterised here on the *completed* threshold, whose edge coincides with the corollary’s radius at the null scorecaster. Past it coverage fails outright, not merely its rate: the published sufficient condition is necessary there. Section 4 sets both claims beside four vocabularies for the same movement.

2 Setup

The producer. At time t a forecaster fit on data before t induces a score s_t , and a threshold q_t is emitted before y_t is revealed. The deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = \mathbf{1}\{s_t(x_t, y_t) > q_t\}$. The target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID [Angelopoulos et al., 2023] manipulates the threshold on the score scale, generalizing the running-error feedback that adaptive conformal inference [Gibbs and Candès, 2021] and its distribution-shift variant DtACI [Gibbs and Candès, 2024] already apply directly to the threshold:

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

Angelopoulos et al.’s iteration (5). Its r_t saturates: condition (4) requires $r_t(x) \geq b$ for $x \geq c h(t)$ and $r_t(x) \leq -b$ for $x \leq -c h(t)$, with h nonnegative, nondecreasing and sublinear. In words: once the accumulated error clears $c h(t)$ in either direction, the correction is at least a full b in the direction that reduces it. With $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$, Proposition 2 gives $|E_T| \leq c h(T) + 1$, read on realised scores; that is a coverage gap of at most $(c h(T) + 1)/T$, which vanishes because h is sublinear.

One scalar, two readouts, two places. Damping a sequentially updated scalar has two standard readouts, both prior work. A quadratic cost gives linear partial adjustment $\tilde{q}_{t+1} = (1 - w)\hat{q}_{t+1}^{\text{raw}} + w\tilde{q}_t$ [Gârleanu and Pedersen, 2013], which Godahewa et al. [2025] publish as post-processing and Binny and Dixit [2025, Eq. 13] apply to a deployed conformal threshold. A proportional cost gives a no-trade band, moving only when the target is more than τ away, then pulling back exactly τ [Constantinides, 1986, Davis and Norman, 1990]. Either readout has two positions in (1), drawn in Figure 1. (A), Placement A, is the one measured here, and it is the one a designer reaches for, because the deployed number is what has to sit still. *The dead band (L_1) is not a stylistic alternative to the partial-adjustment (L_2) form*; it is the family whose failure Section 3 characterises. Neither is safe by construction: L_2 forfeits Proposition 2’s rate while covering, L_1 past the radius loses coverage, and at the legal $\hat{q} \equiv -b/2$ both lose it. So the tight quantity is the radius, not either form.

The quantity, and the names it already has. We report the deployed threshold’s *travel*, $\mathcal{T}_H = \sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new, and the name is a convenience rather than a contribution. Over a predictive distribution $\sum |\Delta q|$ is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls the same movement *jumpiness* [Zsótér et al., 2009], a *convergence index* [Ehret, 2010] and forecast (in)consistency [Pappenberger et al., 2011], all three imported by Van Belle et al. [2026]; applied control calls it IACER [Sarhadi, 2025]; and

dynamic-regret online learning calls it a path length. Section 4 names ten across four vocabularies, so a fresh one is used here, for one reading only.

Harness. The harness fixes $\alpha = 0.1$ and $b = 2$. Its saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$, with $c = 1$ and $h(t) = \log(t + 2)$, meets (4) with equality. The scorecaster is null and legal, reducing (1) to the pure error integration Proposition 2 addresses, so the readout is the only manipulated variable. Neither choice is neutral, and each sits at an edge of what condition (4) and Theorem 1 allow rather than in the middle of it: (4) is a lower bound, and this saturator meets it exactly rather than exceeding it, while Theorem 1 permits scorecasters anywhere within $b/2$ either side of null, predictable but not required to be constant in t . The adversary is specified, not tie-broken. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$, $\varepsilon = 10^{-9}$, against the *deployed* $\tilde{q}_t$: the $\varepsilon \downarrow 0$ reading of “sets the score at the threshold”, firing the strict indicator. Legality is Theorem 1’s $[-b/2, b/2]$ on both sides, so $b/2 + b/2 = b$ is exactly tight. *That specification names one member of a class*: any adapted $s_t \in [-b/2, b/2]$ is legal, the constant $s_t \equiv -b/2$ included. Every number in Table 1 is measured against the one adversary above; Section 3’s two boundaries are quantified over the class, and they differ.

3 The edge of the admissible set, located and proved

The retention result is published, and tested from outside. Li et al. [2025, Corollary 2] prove that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the deployed threshold keeps (1) in error-integration form and *retains* $|E_T| \leq c h(T) + 1$: the radius is theirs. A downstream partial adjustment of weight w leaves it; Table 1 prices that departure, and five of its eleven arms exceed the bound at $T = 10^6$, and for $w = 0.99$ and 0.999 the excursion is *flat in T* while the bound grows like $\log T$. The forfeit is a property of the smoothing *strength*, not Placement A, and it buys something: travel falls from 28,311 to 202. The forfeit’s sign is counter-intuitive, so it is worth saying what does and does not break. A readout on the completed q_{t+1} touches neither r_t nor the integrator’s argument, so it cannot put (4) at risk. What it does instead is delay a correction that has already been computed, so the accumulator excurses *further* rather than less: smoothing the output does not damp E_t but feeds it.

Table 1: Placement A: a one-scalar readout on the *completed* threshold, under one deterministic adversary, on one pass to $T = 10^6$; all eleven arms run are listed. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, and “Sat.” the fraction of rounds on which (4) delivers its full $\pm b$. Every number is measured against the adversary Section 2 specifies.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t \ 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5

The L_1 dead band leaves the radius, and the edge is a theorem. A dead band pulls the deployed threshold back by exactly τ when it releases and holds it otherwise, so $|\hat{q}_{t+1} - q_{t+1}| \leq \tau$ at every round and “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under saturation the deployed value is pinned τ below the integrator’s ceiling, and the deployed value’s realised supremum is $b - \tau$ to the printed digit at every $\tau \geq 0.9$. Outside the radius the answer turns on two different levels of the integrator, and condition (4) names only one of them. The two are the largest correction r_t ever issues and the correction it is guaranteed to issue once the accumulator has cleared the radius; where they differ, a boundary written with the first of them over-certifies. Write $A_t^+ = \sup_x r_t(x)$ and $A_t^- = \inf_x r_t(x)$ for the integrator’s extremes, and $\Lambda_t^+ = \inf_{x \geq c h(t)} r_t(x)$, $\Lambda_t^- = \sup_{x \leq -c h(t)} r_t(x)$ for the levels it is *guaranteed* to reach once the accumulator clears the radius. Condition (4) says exactly $\Lambda_t^+ \geq b$ and $\Lambda_t^- \leq -b$; that $\Lambda_t^\pm = A_t^\pm$, that the integrator reaches its extremes where (4) puts them, is a further hypothesis and not one (4) supplies. Set

$$\tau^{\pm} = \sup_t (|A_t^\pm| \pm \hat{q}_{t+1}) - b/2, \quad \sigma^\pm = \inf_t (|\Lambda_t^\pm| \pm \hat{q}_{t+1}) - b/2. \quad (2)$$

117 **Failure, over the whole admissible class.** If $\tau > \tau^{*+}$ then $\hat{q}_{t+1} + A_t^+ - \tau < b/2$ at every round,
 118 so $\tilde{q}_{t+1} \leq \max(\tilde{q}_t, q_{t+1} - \tau)$ makes $\{\tilde{q} < b/2\}$ absorbing: a release upward lands below $b/2$, and a
 119 hold or a release downward cannot lift the deployed value at all. Against $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$ the
 120 indicator is then $1\{\tilde{q}_t < b/2\}$ and every round misses. If $\tau \geq \tau^{*-}$ the constant legal play $s_t \equiv -b/2$
 121 makes $\{\tilde{q} \geq -b/2\}$ absorbing by the mirror inequality, so no round can miss and the realised rate
 122 is 0 rather than α . Either way long-run coverage is lost, with $|E_T|$ linear in T . Nothing here asks $\hat{q}$
 123 to be constant, asks the extremes to be attained anywhere, or asks $A_t^\pm$ to be finite beyond what the
 124 hypothesis already forces. **Retention, keyed to the reach rather than the supremum.** If $\tau \leq \sigma^+$
 125 then $E_{t-1} \geq ch(t-1)$ forces $q_t \geq b/2 + \tau$ and so $\tilde{q}_t \geq b/2$, whence $\text{err}_t = 0$ against every legal
 126 adversary; if $\tau < \sigma^-$ the mirror forces $\text{err}_t = 1$ below $-ch(t-1)$. The accumulator cannot then
 127 leave $[-ch(t), ch(t)]$ by more than one round, and $|E_T| \leq ch(T) + 1$: Proposition 2's own bound,
 128 retained verbatim rather than weakened to $o(T)$. Using condition (4) alone this already certifies
 129 $\tau \leq b/2 + \inf_t \hat{q}_t$ and $\tau < b/2 - \sup_t \hat{q}_t$, which for a constant scorecaster is Corollary 2's radius at
 130 $\mu_t = 1$.

131 **The two directions meet, with no gap, on the sub-class the five settings inhabit.** If $\hat{q}_t \equiv \hat{q}$, if
 132 $\Lambda_t^\pm = A_t^\pm$, and if neither depends on t , then $\sigma^\pm = \tau^{*\pm}$ and **coverage is retained against every**
 133 **legal adapted adversary if and only if $\tau \leq \tau^{*+}$ and $\tau < \tau^{*-}$.** For a symmetric saturator that is
 134 one line: coverage if and only if $\tau \leq \sup_x r_t(x) - b/2 - |\hat{q}|$, with the endpoint included only when
 135 $\hat{q} < 0$. Every setting run here is in that sub-class, so the sharp form applies to all five. **What was**
 136 **printed before is the failure half.** τ^{*+} is the rule Figures 2 and 3 draw, and all 19 dead-band runs
 137 agree with it read against the adversary the harness plays. It is not a retention condition, and Table 2
 138 (App. A) separates the two with legal objects only. The supremum can sit where the dynamics never
 139 go: $|E_t| \leq (1 - \alpha)t$, so a saturator equal to the harness's inside $|x| \leq t^2$ and to $\pm 10b$ outside it
 140 satisfies (4), reads $\tau^{*+} = 19$, and fails at $\tau = 1.001$. A legal time-varying $\hat{q}$ splits $\sup_t$ from $\inf_t$:
 141 $\hat{q}_t = b/2$ on t a power of two and 0 elsewhere reads $\tau^{*+} = 2$ against $\sigma^+ = 1$. And that printed rule
 142 omits τ^{*-} entirely.




Figure 2: Realised miscoverage against dead-band width τ at the null scorecaster, $\hat{q} \equiv 0$ with a saturator of level b , under the *specified* adversary, where the upward boundary τ^{*+} of (2) is 1. (a) The full sweep: filled markers run to $T = 10^6$ and the four widest bands, drawn open, to $T = 2 \times 10^5$; blue covers, orange forfeits, and the dashed rule is τ^{*+} . (b) The same setting at $T = 10^5$ over eleven measured widths, placed by $\tau - \tau^{*+}$ on a symlog axis whose linear zone is $|\tau - \tau^{*+}| \leq 10^{-3}$, which magnifies the neighbourhood of τ^{*+} : visual spacing here is not proportional to true distance in τ . Unlabelled minor ticks mark $\tau = 0.95$ and 1.05 . The grey band spans the widest covering width to the narrowest failing one; Figure 3 gives the other four $(r_t, \hat{q})$ settings. The same eleven widths under the *mirror* adversary, which locates τ^{*-} , are this run's exact-threshold column read under the strict indicator.

143 **The sharp form is one number, and its hypotheses are what buy it.** Specialise to a symmetric
 144 saturator that meets (4) with equality and attains its extremes there ($\Lambda_t^\pm = A_t^\pm = \pm b$; the harness's
 145 $r_t = b \text{clip}(x/(ch(t)), \pm 1)$ is one, Section 2), and to a constant $\hat{q}$. Then

$$\min(\tau^{*+}, \tau^{*-}) = b/2 - |\hat{q}|, \quad (3)$$

146 with the endpoint admissible only when $\hat{q} < 0$: exactly Corollary 2's radius at $\mu_t = 1$, now as an
 147 if-and-only-if. **The necessity direction above assumes none of it:** $\tau > \tau^{*+}$, or $\tau \geq \tau^{*-}$, forfeits
 148 coverage against every admissible r_t and every legal predictable $\{\hat{q}_t\}$, extremes attained or not, and




Figure 3: All 19 dead-band runs, at the five admissible $(r_t, \hat{q})$ settings: rows are settings ordered by τ^{*+} , and the shared horizontal axis is τ . Blue markers cover at $\approx \alpha$ and orange markers return 1.000000; filled markers run to $T = 10^6$ and open markers to $T = 2 \times 10^5$, all under the *specified* adversary. The rule on each row is that setting's upward boundary $\tau^{*+} = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, printed at the right, and the tinted band is $\tau \leq \tau^{*+}$. Two rules fall off the axis and are drawn as arrows: $\tau^{*+} = 0$ at $\hat{q} \equiv -b/2$, and τ^{*+} unbounded under Angelopoulos et al.'s tangent integrator. The mirror boundary τ^{*-} of (2) is not drawn here and differs from τ^{*+} at the two non-null scorecasters.

that unrestricted statement is what this section claims. The two hypotheses are load-bearing rather than decorative: the unattained supremum and the varying- $\hat{q}$ split two paragraphs above are exactly what they rule out. Equation (3) is where the reach meets the supremum, not evidence that they always do.

The mirror boundary is already measured, on this paper's own grid. The τ -cliff run's exact-threshold column under the strict indicator is the mirror adversary $s_t \equiv -b/2$ (one error process, $\text{err}_t = \mathbf{1}\{\hat{q}_t < -b/2\}$), and on the same eleven widths at the null scorecaster it locates $\tau^{*-} = 1$, closed on the failing side, exactly where (2) puts it. Both 1.000000 and 0.000000 are failures of the same order, $|E_T| = (1 - \alpha)T$ and αT , **so the failure criterion is $|E_T|$, not the miscoverage rate;** and once $\hat{q}$ varies the rate stops being a tell at all.

Both degenerate edges, and what they need. An unbounded supremum is not on its own enough for $\tau^{*+} = \infty$ to mean what it looks like: the same unreachable-set construction with $r_t(x) = x$ past $|x| > t^2$ has no finite supremum and fails at $\tau = 1.5$. What the tangent family has, and that construction does not, is a genuine reach: $\Lambda_t(X) = \tan(\arctan(b)X/(ch(t)))$, so the accumulator level at which the deployed threshold is forced to $b/2$ is $\kappa(\tau)ch(t)$ with $\kappa(\tau) = \arctan(b/2 + \tau - \hat{q})/\arctan(b)$, finite at every finite τ . Coverage is retained at every width, as measured; but $\kappa(\tau) > 1$ once $\tau > b/2 + \hat{q}$, so Proposition 2's own *constant* is forfeited past $\tau = b/2$ while coverage is not. At $\tau = 1.5$, $T = 10^6$ the excursion is 15.10 against Proposition 2's 14.8155 and against $\kappa(1.5)ch(T) + 1 = 15.8530$. At the other edge, $\sup_t \hat{q}_t = -b/2$ forces $\hat{q} \equiv -b/2$, since every $\hat{q}_t$ is at least $-b/2$ already. There $q_{t+1} - \tau \leq -b/2 + b - \tau < b/2$ for every $\tau > 0$ and every t , and $\tilde{q}_1 = 0 < b/2$, so the failing set is absorbing from round 1 with no transient: the admissible window is the single point $\tau = 0$, and $\tau = 0.5$ returns 1.000000, as does partial adjustment at $w = 0.999$.

What that establishes is that the published condition is tight. At $\mu_t = 1$, $\hat{q} \equiv 0$ the boundary is exactly Corollary 2's radius, crossing it costs coverage rather than the rate ($\tau = 0.9$ covers at 0.100012, $\tau = 1.5$ returns 1.000000 with $\max_t |E_t| = 900,000$, 60,747 times Proposition 2's bound at $T = 10^6$; Table 1), and (2) makes that an if-and-only-if against the whole legal adversary class rather than a pattern across five settings. **The claim is that necessity: a tightness result about someone else's sufficient condition**, smaller and more precise than a forfeit of the rate by post-processing. Elsewhere it stays sufficient only, and now provably so: under the tangent integrator Corollary 2 permits $\tau \leq b/2$ while no finite width forfeits coverage against any legal adversary, because the reach is genuine at every level. At $\hat{q} \equiv +b/2$ the looseness goes the other way and disappears: Corollary 2 permits radius 0, and (2) permits no width at all against the mirror adversary, the unsmoothed control included, though $\tau = 1.5$ covers there against the specified one. The condition therefore binds at some legal settings and is slack at others, and that asymmetry is what makes "tight" a located claim rather than a global one: what is stated here is where the radius is sharp, and equally where it is not.

185 **The constructive reading, and it is the corollary’s rather than ours.** Placement B keeps $\{\hat{q}_t\}$
 186 inside the admissible radius, so Theorem 1 carries it already. What Section 3 adds is that the band
 187 has two edges, both computable before deployment from the integrator’s guaranteed reach and the
 188 scorecaster’s range, that crossing either is not a graceful degradation, and that between $\sigma^\pm$ and $\tau^{*\pm}$,
 189 which separate only when $\hat{q}$ varies in t , neither direction is settled here.

190 **The boundary read as a bet.** The boundary is stated in $|E_T|$, which is the net gain of a Skeptic
 191 who pays α each round for err_t , so it has a betting reading. Impose, for that reading alone and on
 192 nothing else here, the null that $\text{err}_t \mid \mathcal{F}_{t-1}$ is Bernoulli(α), with $\mathcal{F}_{t-1}$ the history through round
 193 $t - 1$: past the edge an anytime-valid test of the deployed intervals’ calibration convicts the $\tau = 1.5$
 194 arm on its fifth interval at the 5% level, and inside the edge the same test cannot reject at that level at
 195 any horizon. It adds a diagnostic and a reading, not a result: the test is Ville’s inequality applied to
 196 the standard martingale $M_t = \prod_{i \leq t} (1 + \lambda(\text{err}_i - \alpha))$, whose logarithm depends on the path only
 197 through t and E_t [Shafer and Vovk, 2019, Ramdas et al., 2023, Vovk and Wang, 2021]; λ is fixed in
 198 $[-1/(1 - \alpha), 1/\alpha]$, and its sign above is chosen after the fact, not pre-registered; M_t is an e-value
 199 for that imposed null alone, the accumulator itself is signed and is neither an e-value nor an e-process,
 200 and no result above assumes that null or is certified by that test.

201 4 Where this sits

202 **One quantity, four names.** Four literatures name the movement of a sequentially updated fore-
 203 cast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC,
 204 MQC/SMQC, congruence, nervousness); *jumpiness*, *convergence index*, *(in)consistency*; and *switch-*
 205 *ing cost*. Online conformal prediction gives distribution-free long-run coverage for any bounded
 206 scorecaster [Angelopoulos et al., 2023, Thm 1, Prop. 2]. Forecast stability treats the movement as
 207 a functional, read out as post-processing [Godahewa et al., 2025, Eq. 3], against a priced decision
 208 to switch [Genov et al., 2026, §4.2], [Van Belle et al., 2026]. Verification work uses a matched-
 209 level design against real producers from 2012 [Pinson and Girard, 2012, Worsnop et al., 2018].
 210 Switching-cost online learning gives worst-case theory for the penalised problem [Andrew et al.,
 211 2013, Thm 2].

212 The object that settles a deployed interval’s validity (the admissible set around the deployed threshold)
 213 sits in the first of them. Those four vocabularies are where this paper’s claim and its concessions have
 214 to be placed. **What is conceded, by name.** The placement is prior work: Angelopoulos et al. [2023,
 215 Thm. 1, Prop. 2] state coverage, with a rate, for any admissible integrator and any scorecaster, and run
 216 a smoothing-family model in that slot themselves; this paper proves that admissible set tight. Dupuy
 217 et al. [2026, Appendix A] give the generic argument, and Duerst et al. [2024] adopt that scorecaster.
 218 The derivation is prior work too, and this concession matters most. Li et al. [2025, Corollary 2]
 219 prove the retention statement in Conformal PID’s own notation, with Wang and Hyndman [2024] an
 220 independent second instance. Section 2 concedes the smoother object, the metric arithmetic and both
 221 readout maps. Zheng et al. [2025], Wu et al. [2025], Chen et al. [2023], Lade et al. [2026] each place a
 222 smoother inside their own validity argument. One vocabulary out, this is integrator anti-windup with a
 223 feedforward path: free-until-the-bound-binds is its defining specification, not a finding [Galeani et al.,
 224 2009]. Nearest in the other direction, Gao et al. [2025] run the same recursion but gate the *model*
 225 update on the conformal score. Their one non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$
 226 for a sublinear h , supported with a plot rather than derived. Verification’s matched-level design tests
 227 real producers from 2012, with a control stronger than the pair: the whole marginal [Pinson and
 228 Girard, 2012, Worsnop et al., 2018]. Switching-cost online learning adds a Thm. 7 trade-off in which
 229 sublinear regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013, Thm. 7].

230 5 Limitations

231 The result above is narrow in several specific ways. They are listed here rather than left to be found,
 232 and the run-in bold marks where one concession ends and the next begins.

233 **The retention direction is proved with no gap only where the two boundaries meet:** a constant
 234 scorecaster and a saturator attaining its extremes at the condition-(4) radius, which is every setting run
 235 here; a genuinely time-varying $\hat{q}$ leaves the band $\sigma^+ < \tau \leq \tau^{*+}$ open in neither direction (Section 3).
 236 And the excursion is specific to the saturator’s *level*, which condition (4) bounds only from below.

237 **The inherited guarantee is weak where it applies.** Under the tangent integrator $h(t) = t/\log t$,
238 the certified gap $(ch(T) + 1)/T$ decays as $O(1/\log T)$ rather than $O(1/T)$. What is forfeited was
239 not a tight certificate. **The dead-band widths are absolute, and the failure is total only against**
240 **the adversary.** τ is in score units with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$. Under i.i.d. scores it
241 returns 0.249376 at $T = 10^6$, an adversary-free confirmation of the mechanism rather than a retreat.
242 The threshold pins at $b - \tau = 0.5$, and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly.
243 What follows concerns the evidence rather than the statement. **The harness was rebuilt against num-**
244 **bers already read**, which carries less weight than a rerun done without sight of them. **Placement B**
245 **is conceded, not tested:** the argument that it stays inside the admissible radius is Theorem 1's, and no
246 arm run here exercises it. **The i.i.d. check above used one seed** (20260820, $T = 10^6$); **the primary,**
247 **adversarial harness uses none, and only its null-scorecaster edge is bracketed;** the others are
248 one-sided.
249 **What remains open is the band between the two boundaries.** For a genuinely time-varying $\hat{q}$,
250 coverage in $\sigma^+ < \tau \leq \tau^{*+}$ is neither proved nor falsified; Table 2's power-of-two rows sit there, and
251 closing that band is future work, not a sweep this paper has already run.

252 References

- 253 Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan
254 Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness
255 and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages
256 741–763, 2013.
- 257 Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction.
258 *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023. doi: 10.1561/2200000101.
- 259 Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control
260 for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS*
261 *2023)*, 2023. doi: 10.52202/075280-1000.
- 262 Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? Conformal prediction for
263 interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- 264 Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng
265 Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference
266 for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge*
267 *Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- 268 George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political*
269 *Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.
- 270 M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of*
271 *Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.
- 272 Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic
273 forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical
274 Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- 275 Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-
276 aware thresholding in online conformal prediction for time series. In *Artificial Intelligence*
277 *for Healthcare, and Hybrid Models for Coupling Deductive and Inductive Reasoning: First*
278 *International Joint Conference, HC@AIxIA+HYDRA 2025, Bologna, Italy, October 25–26, 2025,*
279 *Proceedings*, volume 2830 of *Communications in Computer and Information Science*, pages
280 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- 281 Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational
282 rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/
283 2010/0480.
- 284 Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-
285 windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.

286 Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep
287 Koopman linear embeddings. *arXiv:2511.12760v2*, 2025.

288 Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction
289 costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.

290 Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Sal-
291 vador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online
292 optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026.
293 doi: 10.1016/j.eswa.2025.129305.

294 Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In
295 *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021.

296 Isaac Gibbs and Emmanuel J. Candès. Conformal inference for online prediction with arbitrary
297 distribution shifts. *Journal of Machine Learning Research*, 25(162):1–36, 2024.

298 Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration
299 and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(2):
300 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.

301 Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador
302 García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4):
303 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.

304 Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal
305 prediction. *arXiv:2512.07770v2*, 2026.

306 Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for
307 multi-horizon time series forecasting. *arXiv:2604.13253v1*, 2026.

308 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
309 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
310 2025. doi: 10.1609/aaai.v39i17.34029.

311 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
312 tion for multi-step forecasting. *arXiv:2508.13362v3*, 2025.

313 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
314 metric in conformal prediction: When shorter intervals are not better. *arXiv:2601.21455v2*, 2026.

315 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
316 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth
317 System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.

318 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
319 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.

320 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
321 and inventory decisions. SSRN working paper 4711817, 2024.

322 Aaditya Ramdas, Peter Grünwald, Vladimir Vovk, and Glenn Shafer. Game-theoretic statistics and
323 safe anytime-valid inference. *Statistical Science*, 38(4):576–601, 2023. doi: 10.1214/23-STS894.

324 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
325 machine learning controller? *arXiv:2503.14379v3*, 2025.

326 Glenn Shafer and Vladimir Vovk. *Game-Theoretic Foundations for Probability and Finance*. Wi-
327 ley Series in Probability and Statistics. Wiley, 2019. ISBN 9780470903056. doi: 10.1002/
328 9781118548035.

329 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
330 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
331 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.

- 332 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
333 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
334 007.
- 335 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
336 probabilistic forecasts. arXiv:2605.28531v1, 2026.
- 337 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
338 arXiv:2607.26577v1, 2026.
- 339 Vladimir Vovk and Ruodu Wang. E-values: Calibration, combination and applications. *The Annals*
340 *of Statistics*, 49(3):1736–1754, 2021. doi: 10.1214/20-AOS2020.
- 341 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
342 ing. arXiv:2410.13115v2, 2024.
- 343 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
344 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
345 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.
- 346 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
347 inference for time series. arXiv:2502.00818v2, 2025.
- 348 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
349 forecasting systems. arXiv:2506.05776v3, 2026.
- 350 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
351 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.
- 352 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
353 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
354 doi: 10.1175/2009MWR2960.1.

355 Code and data

356 The simulator, its frozen configuration, the results/ files behind every figure and table, the gen-
357 erators that read them, and this project’s audit trail across sessions are at [https:// \[REPOSITORY](https://[REPOSITORY LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION])
358 [LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION\]](https://[REPOSITORY LINK - INSERTED BY AUTHOR PRIOR TO SUBMISSION]).

359 A Supporting measurements

Table 2: Legal $(r_t, \hat{q})$ configurations on which the printed τ^* over-certifies, each under two legal adversaries. “Proved window” is where (2) certifies retention against every legal adversary; between it and the failure boundaries $\tau^{*\pm}$ neither direction is claimed, and the rows built on a supremum the dynamics cannot reach or never attain sit there. Here $\emptyset$ marks a configuration for which no width is certified, $\tau = 0$ included. $T = 10^5$ except the two power-of-two rows, at $T = 10^6$. The last row carries no dead band at all.

$(r_t, \hat{q})$ configuration	τ	printed τ^*	proved window	$s_t = \hat{q}_t + \varepsilon$		$s_t \equiv -b/2$	
				miscov.	$\max_t E_t $	miscov.	$\max_t E_t $
$\hat{q} \equiv +0.4$, level b	0.599	1.4	$\tau < 0.6$	0.100010	7.20	0.099890	11.60
$\hat{q} \equiv +0.4$, level b	0.8	1.4	$\tau < 0.6$	0.100050	8.70	0.000000	10,000
$\sup_x r_t = 10b$, off the reach	1.001	19	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t$ unbounded, off the reach	1.5	∞	$\tau < 1$	1.000000	90,000	0.000000	10,000
$A^+ = 4b$, $A^- = -b$	1.5	7	$\tau < 1$	0.100020	4.30	0.000000	10,000
level $4b$, $\hat{q} \equiv 0$	6.999	7	$\tau < 7$	0.100070	12.20	0.100050	11.60
level $4b$, $\hat{q} \equiv 0$	7.0	7	$\tau < 7$	0.100090	12.20	0.000000	10,000
$\sup_x r_t = 2b$, unattained	2.99	3	$\tau < 1$	0.100320	211.10	0.098010	210.50
$\sup_x r_t = 2b$, unattained	3.0	3	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t = 2b$, attained	3.0	3	$\tau < 3$	0.100000	12.20	0.000000	10,000
$\hat{q}_t = b/2$ on powers of 2	1.0	2	$\emptyset$	0.100009	14.30	0.000000	100,000
$\hat{q}_t = b/2$ on powers of 2	1.9	2	$\emptyset$	0.110960	84,745.70	0.000000	100,000
$\hat{q} \equiv +b/2$, no band	0	2	$\emptyset$	0.100000	0.90	0.000000	10,000