
Where the Admissible Radius Binds: a Correction and a Measured Boundary in Online Conformal Prediction

Palaash Gang
Independent Researcher
palaashgang@gmail.com

Abstract

1 A printed claim that the scorecaster breaks online conformal coverage meets its
2 refutation by overlapping authors; this paper joins them and locates the edge
3 of the set that refutation relies on. Conformal PID states long-run coverage for
4 any bounded scorecaster; a corollary keeps it for predictable perturbations of the
5 deployed threshold inside an admissible radius. The claim was acted on; placement
6 and derivation are conceded by name. At the null scorecaster that condition is *tight*.
7 A legal L_1 dead band on the *completed* threshold leaves that set at its radius, and
8 past it miscoverage is 1.000000 against $\alpha = 0.1$. Coverage goes, not only the rate.
9 The edge belongs to the saturator–scorecaster pair, not the readout; elsewhere the
10 condition is sufficient only. Under the equally legal $\hat{q} \equiv +b/2$ the failing band
11 covers; at $\hat{q} \equiv -b/2$ both it and partial adjustment at $w = 0.999$ miscover.

1 Introduction

13 A refereed paper states that a scorecaster breaks Conformal PID’s coverage guarantee. A later one,
14 by the same first and last authors, proves in the same notation that a bounded perturbation of the
15 deployed threshold does not. The later cites the earlier, so this is ordinary self-correction, and the
16 joining is what this paper adds.

17 An online conformal method [Angelopoulos and Bates, 2021] emits a threshold before each obser-
18 vation and moves it. Forecast stability names, scores and prices that movement [Tunc et al., 2013,
19 Godahewa et al., 2025, Van Belle et al., 2023, Pritularga and Kourentzes, 2024, Genov et al., 2026,
20 Van Belle et al., 2026]. Varying temporal structure at fixed level came earlier, in forecast verification,
21 matching predictive marginals on real producers [Gneiting et al., 2007, Pinson and Girard, 2012,
22 Worsnop et al., 2018]. Coverage and mean length under-describe a deployed interval [Min et al.,
23 2026, Vaze, 2026]. None of it is claimed here.

24 **The record carries a claim and its refutation; this paper joins them.** Conformal PID [An-
25 gelopoulos et al., 2023] adds a *scorecaster* $\hat{q}_{t+1}$ to a saturating error integrator, asking only that it be
26 predictable and lie in $[-b/2, b/2]$: “any function of the past” (Theorem 1). Such a readout is already
27 quantified over. The claim reads “[s]ince using a scorecaster (D-part) in C-PID breaks the theoretical
28 coverage guarantee” [Li and Rodríguez, 2025], stated and acted on in its baselines paragraph, ver-
29 batim on refereed printed page 18443. It is refuted twice, with proof. Li et al. [2025, Corollary 2]
30 add any predictable d_t , $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ with $\mu_t \in [0, 1]$ predictable, to the deployed threshold;
31 since $\hat{q}_t + d_t$ still lies in $[-b/2, b/2]$ it is itself a legal scorecaster, so the result “has exactly the CPID
32 error-integration form” and $|E_T| = o(T)$. The same generic-slot argument carries AcMCP’s long-run
33 coverage [Wang and Hyndman, 2024]. Both identifiers and “scorecaster” were searched across arXiv
34 full text, arXiv abstracts, OpenReview, and Semantic Scholar’s six citers of Li and Rodríguez [2025]:
35 none returns both the claim and the corollary. Hu et al. [2026] still call scorecaster choice “arbitrary”
36 and lacking “principled guidance” while asserting valid coverage: model selection, not validity. **We**
37 **claim neither the placement nor the derivation.**




Figure 1: One scalar readout (grey) in its two places in (1); dashed is the loop r_t closes. In (B) the loop closes on the integrator’s own output, and both readouts keep $\{\hat{q}_t\}$ in the convex hull of $\hat{q}_1$ and the raw scorecaster’s range, inside Corollary 2’s admissible set. In (A) it closes on $\tilde{q}_{t+1}$, which it never computed, and the deployed sequence leaves that set.

38 **Where that refutation stops, the condition is tight.** Corollary 2 is a *retention* result, silent outside
 39 its radius, and the question it leaves is whether that radius is sharp. Section 3 exhibits a legal
 40 perturbation leaving it: the L_1 dead-band family characterised here on the *completed* threshold,
 41 whose edge is the corollary’s radius at the null scorecaster. Past it coverage goes, not only the
 42 rate: the published sufficient condition is necessary there. Appendix C sets both claims beside four
 43 vocabularies for the same movement, and reads E_t as a game-theoretic capital process.

44 2 Setup

45 **The producer.** A forecaster fit before t induces a score s_t , and a threshold q_t is emitted before
 46 y_t is revealed. The deployed set is $C_t = \{y : s_t(x_t, y) \leq q_t\}$ and $\text{err}_t = \mathbf{1}\{s_t(x_t, y_t) > q_t\}$. The
 47 target is long-run coverage, $T^{-1} \sum_{t \leq T} \text{err}_t = \alpha + o(1)$, with no model on the data. Conformal PID
 48 [Angelopoulos et al., 2023] manipulates the threshold on the score scale, generalizing the running-
 49 error feedback that adaptive conformal inference [Gibbs and Candès, 2021] and its distribution-shift
 50 variant DtACI [Gibbs and Candès, 2024] already apply directly to the threshold:

$$q_{t+1} = \hat{q}_{t+1} + r_t \left(\sum_{i \leq t} (\text{err}_i - \alpha) \right), \quad (1)$$

51 Angelopoulos et al.’s iteration (5). Its r_t saturates: condition (4) requires $|r_t(x)| \geq b$ with the sign of
 52 x once $|x| \geq c h(t)$, for h nonnegative, nondecreasing and sublinear. With $E_t = \sum_{i \leq t} (\text{err}_i - \alpha)$,
 53 Proposition 2 gives $|E_T| \leq c h(T) + 1$, read on realised scores.

54 **One scalar, two readouts, two places.** Damping a sequentially updated scalar has two standard
 55 readouts, both prior work. A quadratic cost gives linear partial adjustment $\hat{q}_{t+1} = (1 - \lambda)\hat{q}_{t+1}^{\text{raw}} + \lambda\hat{q}_t$
 56 [Gârleanu and Pedersen, 2013], published as post-processing [Godahewa et al., 2025] and applied
 57 to a deployed threshold [Binny and Dixit, 2025, Eq. 13]. A proportional cost gives a no-trade band,
 58 moving only when the target is more than τ away, then pulling back exactly τ [Constantinides,
 59 1986, Davis and Norman, 1990]. Either readout has two positions in (1), drawn in Figure 1; (A)
 60 is measured here. *The dead band (L_1) is not a stylistic alternative to the partial-adjustment (L_2) form;*
 61 *it is the family whose failure Section 3 characterises. Neither is safe by construction: L_2 forfeits*
 62 *Proposition 2’s rate while covering, L_1 past the radius loses coverage, and at the legal $\hat{q} \equiv -b/2$ both*
 63 *lose it. What is tight is the radius, not a form.*

64 **The quantity, and the names it already has.** We report the deployed threshold’s *travel*, $\mathcal{T}_H =$
 65 $\sum_{t=2}^H |q_t - q_{t-1}|$, after actuator travel in process control. The arithmetic is not new. Over a predictive
 66 distribution $\sum |\Delta q|$ is the Multi-Quantile Change [Zanotti, 2026, v3, §3.4.1]; verification calls it
 67 *jumpiness*, a *convergence index*, (in)consistency [Zsótér et al., 2009, Ehret, 2010, Pappenberger
 68 et al., 2011]; applied control, IACER [Sarhadi, 2025]; dynamic-regret online learning, a path length.
 69 Table 3’s caption names ten, so a fresh one is used here.

70 **Harness.** The harness fixes $\alpha = 0.1$ and $b = 2$. Its saturator $r_t(x) = b \text{clip}(x/(c h(t)), -1, 1)$,
 71 with $c = 1$ and $h(t) = \log(t + 2)$, meets (4) with equality. The scorecaster is null and legal, reducing
 72 (1) to the pure error integration Proposition 2 addresses. Neither choice is neutral: (4) is a lower
 73 bound met exactly here, and Theorem 1 permits scorecasters $b/2$ either side of null, predictable
 74 but not constant in t . The adversary is specified, not tie-broken. It plays $s_t = \text{clip}(\hat{q}_t + \varepsilon, \pm b/2)$,
 75 $\varepsilon = 10^{-9}$, against the *deployed* $\hat{q}_t$: the $\varepsilon \downarrow 0$ reading of “sets the score at the threshold”, firing

76 the strict indicator. Legality is Theorem 1’s $[-b/2, b/2]$ on both sides, so $b/2 + b/2 = b$ is exactly
 77 tight. *That specification names one member of a class:* any adapted $s_t \in [-b/2, b/2]$ is legal, the
 78 constant $s_t \equiv -b/2$ included. Every number in Table 1 is measured against the one adversary above;
 79 Section 3’s two boundaries are quantified over the class, and they differ.

80 3 The edge of the admissible set, located and proved

81 **The retention result is published, and tested from outside.** Li et al. [2025, Corollary 2] prove
 82 that a predictable d_t with $|d_t| \leq \mu_t(b/2 - |\hat{q}_t|)$ on the deployed threshold keeps (1) in error-integration
 83 form and *retains* $|E_T| \leq c h(T) + 1$: the radius is theirs. A downstream partial adjustment of weight
 84 w leaves it; Table 1 prices that departure, and five of its eleven arms exceed the bound at $T = 10^6$,
 85 and for the strongest two the excursion is *flat in T* while the bound grows like $\log T$. The forfeit is a
 86 property of the smoothing *strength*, not Placement A, and it buys something: travel falls from 28,311
 87 to 202. The forfeit’s sign is counter-intuitive. A readout on the completed q_{t+1} touches neither r_t nor
 88 the integrator’s argument, so it cannot put (4) at risk. It delays a correction already computed, and the
 89 accumulator excurses *further*. Smoothing the output does not damp E_t ; it feeds it.

90 **The L_1 dead band leaves the radius, and the edge is a theorem.** A dead band pulls the deployed
 91 threshold back by exactly τ when it releases and holds it otherwise, so $|\tilde{q}_{t+1} - q_{t+1}| \leq \tau$ at every
 92 round and “inside Corollary 2” reads $\tau \leq \mu_t(b/2 - |\hat{q}_t|)$. Under saturation the deployed value
 93 is pinned τ below the integrator’s ceiling, and its realised supremum is $b - \tau$ to the printed digit
 94 at every $\tau \geq 0.9$. Outside the radius two levels decide the answer and condition (4) names only
 95 one of them. Write $A_t^+ = \sup_x r_t(x)$ and $A_t^- = \inf_x r_t(x)$ for the integrator’s extremes, and
 96 $\Lambda_t^+ = \inf_{x \geq c h(t)} r_t(x)$, $\Lambda_t^- = \sup_{x \leq -c h(t)} r_t(x)$ for the levels it is *guaranteed* to reach once the
 97 accumulator clears the radius. Condition (4) says exactly $\Lambda_t^+ \geq b$ and $\Lambda_t^- \leq -b$; that $\Lambda_t^\pm = A_t^\pm$,
 98 that the integrator reaches its extremes where (4) puts them, is a further hypothesis and not one (4)
 99 supplies. Set

$$\tau^{*\pm} = \sup_t (|A_t^\pm| \pm \hat{q}_{t+1}) - b/2, \quad \sigma^\pm = \inf_t (|\Lambda_t^\pm| \pm \hat{q}_{t+1}) - b/2. \quad (2)$$

100 **Failure, over the whole admissible class.** If $\tau > \tau^{*+}$ then $\hat{q}_{t+1} + A_t^+ - \tau < b/2$ at every round,
 101 so $\tilde{q}_{t+1} \leq \max(\tilde{q}_t, q_{t+1} - \tau)$ makes $\{\tilde{q} < b/2\}$ absorbing: a release upward lands below $b/2$, and a
 102 hold or a release downward cannot lift the deployed value at all. Against $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$ the
 103 indicator is then $\mathbf{1}\{\tilde{q}_t < b/2\}$ and every round misses. If $\tau \geq \tau^{*-}$ the constant legal play $s_t \equiv -b/2$
 104 makes $\{\tilde{q} \geq -b/2\}$ absorbing by the mirror inequality, so no round can miss and the realised rate
 105 is 0 rather than α . Either way long-run coverage is lost, with $|E_T|$ linear in T . Nothing here asks $\hat{q}$
 106 to be constant, asks the extremes to be attained anywhere, or asks $A_t^\pm$ to be finite beyond what the
 107 hypothesis already forces. **Retention, keyed to the reach rather than the supremum.** If $\tau \leq \sigma^+$
 108 then $E_{t-1} \geq c h(t-1)$ forces $q_t \geq b/2 + \tau$ and so $\tilde{q}_t \geq b/2$, whence $\text{err}_t = 0$ against *every* legal
 109 adversary; if $\tau < \sigma^-$ the mirror forces $\text{err}_t = 1$ below $-c h(t-1)$. The accumulator cannot then
 110 leave $[-c h(t), c h(t)]$ by more than one round, and $|E_T| \leq c h(T) + 1$: Proposition 2’s own bound,
 111 retained verbatim rather than weakened to $o(T)$. Using condition (4) alone this already certifies
 112 $\tau \leq b/2 + \inf_t \hat{q}_t$ and $\tau < b/2 - \sup_t \hat{q}_t$, which for a constant scorecaster is Corollary 2’s radius at
 113 $\mu_t = 1$.

114 **The two directions meet, with no gap, on the sub-class the five settings inhabit.** If $\hat{q}_t \equiv \hat{q}$, if
 115 $\Lambda_t^\pm = A_t^\pm$, and if neither depends on t , then $\sigma^\pm = \tau^{*\pm}$ and **coverage is retained against every legal**
 116 **adapted adversary if and only if $\tau \leq \tau^{*+}$ and $\tau < \tau^{*-}$.** For a symmetric saturator that is one line:
 117 coverage if and only if $\tau \leq \sup_x r_t(x) - b/2 - |\hat{q}|$, with the endpoint included only when $\hat{q} < 0$.
 118 Every setting run here is in that sub-class, so the sharp form applies to all five. **What was printed**
 119 **before is the failure half.** τ^{*+} is the rule Figures 2 and 3 (App. A) draw, and all 19 dead-band runs
 120 agree with it read against the adversary the harness plays. It is not a retention condition, and Table 2
 121 (App. A) separates the two with legal objects only. The supremum can sit where the dynamics never
 122 go: $|E_t| \leq (1 - \alpha)t$, so a saturator equal to the harness’s inside $|x| \leq t^2$ and to $\pm 10b$ outside it
 123 satisfies (4), reads $\tau^{*+} = 19$, and fails at $\tau = 1.001$. A legal time-varying $\hat{q}$ splits $\sup_t$ from $\inf_t$:
 124 $\hat{q}_t = b/2$ on t a power of two and 0 elsewhere reads $\tau^{*+} = 2$ against $\sigma^+ = 1$. And τ^{*-} was missing
 125 entirely.

The mirror boundary is already measured, on this paper’s own grid. Under the strict indicator both $s_t \equiv -b/2$ and $s_t = \text{clip}(\hat{q}_t, \pm b/2)$ give $\text{err}_t = \mathbf{1}\{\hat{q}_t < -b/2\}$, so they are one error process. The τ -cliff run’s exact-threshold column under that indicator is therefore the mirror adversary, measured at the null scorecaster on the same eleven widths, and it locates $\tau^{*-} = 1$, closed on the failing side, exactly where (2) puts it; Appendix A prints the eleven readings. **So the failure criterion is $|E_T|$, not the miscoverage rate.** Both 1.000000 and 0.000000 are failures of the same order, $|E_T| = (1 - \alpha)T$ and αT ; and once $\hat{q}$ varies the rate stops being a tell at all. At $\tau = 1.9$ with the power-of-two scorecaster the rate is 0.110960 against $\alpha = 0.1$ while $\max_t |E_t|$ is 84,745.70 at $T = 10^6$, so a reader watching only the rate would call that coverage.

Both degenerate edges, and what they need. An unbounded supremum is not on its own enough for $\tau^{*+} = \infty$ to mean what it looks like: the same unreachable-set construction with $r_t(x) = x$ past $|x| > t^2$ has no finite supremum and fails at $\tau = 1.5$. What the tangent family has, and that construction does not, is a genuine reach: $\Lambda_t(X) = \tan(\arctan(b)X/(ch(t)))$, so the accumulator level at which the deployed threshold is forced to $b/2$ is $\kappa(\tau)ch(t)$ with $\kappa(\tau) = \arctan(b/2 + \tau - \hat{q})/\arctan(b)$, finite at every finite τ . Coverage is retained at every width, as measured; but $\kappa(\tau) > 1$ once $\tau > b/2 + \hat{q}$, so Proposition 2’s own *constant* is forfeited past $\tau = b/2$ while coverage is not. At $\tau = 1.5$, $T = 10^6$ the excursion is 15.10 against Proposition 2’s 14.8155 and against $\kappa(1.5)ch(T) + 1 = 15.8530$. At the other edge, $\sup_t \hat{q}_t = -b/2$ forces $\hat{q} \equiv -b/2$, since every $\hat{q}_t$ is at least $-b/2$ already. There $q_{t+1} - \tau \leq -b/2 + b - \tau < b/2$ for every $\tau > 0$ and every t , and $\tilde{q}_1 = 0 < b/2$, so the failing set is absorbing from round 1 with no transient: the admissible window is the single point $\tau = 0$, and $\tau = 0.5$ returns 1.000000, as does partial adjustment at $w = 0.999$.

What that establishes is that the published condition is tight. At $\mu_t = 1$, $\hat{q} \equiv 0$ the boundary is exactly Corollary 2’s radius, crossing it costs coverage rather than the rate, and (2) makes that an if-and-only-if against the whole legal adversary class rather than a pattern across five settings. **The claim is that necessity: a tightness result about someone else’s sufficient condition,** smaller and more precise than a forfeit of the rate by post-processing. Elsewhere it stays sufficient only, and now provably so: under the tangent integrator Corollary 2 permits $\tau \leq b/2$ while no finite width forfeits coverage against any legal adversary, because the reach is genuine at every level. At $\hat{q} \equiv +b/2$ the looseness goes the other way and disappears: Corollary 2 permits radius 0, and (2) permits no width at all, the unsmoothed control included. That asymmetry makes “tight” a located claim.

The constructive reading, and it is the corollary’s rather than ours. Placement B keeps $\{\hat{q}_t\}$ inside the admissible radius, so Theorem 1 carries it already. What Section 3 adds is that the band has two edges, both computable before deployment from the integrator’s guaranteed reach and the scorecaster’s range, that crossing either is not a graceful degradation, and that between $\sigma^\pm$ and $\tau^{*\pm}$, which separate only when $\hat{q}$ varies in t , neither direction is settled here.

4 Limitations

The retention direction is proved with no gap only where the two boundaries meet: a constant scorecaster and a saturator attaining its extremes at the condition-(4) radius, which is every setting run here; a genuinely time-varying $\hat{q}$ leaves the band $\sigma^+ < \tau \leq \tau^{*+}$ open in neither direction (Section 3). And $0.63/(1 - w)$ is specific to the saturator’s *level*, which condition (4) bounds only from below. **The inherited guarantee is weak where it applies.** Under the tangent integrator $h(t) = t/\log t$, the certified gap $(ch(T) + 1)/T$ decays as $O(1/\log T)$ rather than $O(1/T)$. What is forfeited was not a tight certificate. **The dead-band widths are absolute, and the failure is total only against the adversary.** τ is in score units with $b = 2$, so the failing $\tau = 1.5$ is $0.75b$. Under i.i.d. scores it returns 0.249376 at $T = 10^6$, an adversary-free confirmation of the mechanism rather than a retreat. The threshold pins at $b - \tau = 0.5$, and uniform scores on $[-b/2, b/2]$ put $P(s > 0.5) = 0.25$ exactly. **The harness was rebuilt against numbers already read,** which carries less weight than a rerun done without sight of them. **Placement B is conceded, not tested. The i.i.d. check above used one seed (20260820, $T = 10^6$); the primary, adversarial harness uses none, and only its null-scorecaster edge is bracketed;** the others are one-sided. **What remains open is the band between the two boundaries.** For a genuinely time-varying $\hat{q}$, coverage in $\sigma^+ < \tau \leq \tau^{*+}$ is neither proved nor falsified; Table 2’s power-of-two rows sit there, and closing that band is future work, not a sweep this paper has already run.

References

- Lachlan L. H. Andrew, Siddharth Barman, Katrina Ligett, Minghong Lin, Adam Meyerson, Alan Roytman, and Adam Wierman. A tale of two metrics: Simultaneous bounds on competitiveness and regret. In *Proceedings of the 26th Conference on Learning Theory (COLT 2013)*, pages 741–763, 2013.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. arXiv:2107.07511v6, 2021.
- Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control for time series prediction. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. doi: 10.52202/075280-1000.
- Allen Emmanuel Binny and Anushri Dixit. Who moved my distribution? Conformal prediction for interactive multi-agent systems. arXiv:2511.11567v1, 2025.
- Jiadong Chen, Yang Luo, Xiuqi Huang, Fuxin Jiang, Yangguang Shi, Tieying Zhang, and Xiaofeng Gao. IPOC: An adaptive interval prediction model based on online chasing and conformal inference for large-scale systems. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 202–212, 2023. doi: 10.1145/3580305.3599396.
- George M. Constantinides. Capital market equilibrium with transaction costs. *Journal of Political Economy*, 94(4):842–862, 1986. doi: 10.1086/261410.
- M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.
- Ricarda Duerst, Jonas Schöley, Julia Hellstrand, and Mikko Myrskylä. Calibrating probabilistic forecast paths on past forecast errors: an application to the Finnish total fertility rate. Technical Report WP-2024-016, Max Planck Institute for Demographic Research, 7 2024.
- Théo Dupuy, Binbin Xu, Stéphane Perrey, Jacky Montmain, and Abdelhak Imoussaten. Relevance-aware thresholding in online conformal prediction for time series. In *Communications in Computer and Information Science*, pages 196–217, 2026. doi: 10.1007/978-3-032-16708-8_17.
- Uwe Ehret. Convergence index: a new performance measure for the temporal stability of operational rainfall forecasts. *Meteorologische Zeitschrift*, 19(5):441–451, 2010. doi: 10.1127/0941-2948/2010/0480.
- Sergio Galeani, Sophie Tarbouriech, Matthew Turner, and Luca Zaccarian. A tutorial on modern anti-windup design. *European Journal of Control*, 15(3-4):418–440, 2009. doi: 10.3166/ejc.15.418-440.
- Ben Gao, Jordan Patracone, Stéphane Chrétien, and Olivier Alata. Conformal online learning of deep Koopman linear embeddings. arXiv:2511.12760v2, 2025.
- Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.
- Evgenii Genov, Julian Ruddick, Christoph Bergmeir, Majid Vafaeipour, Thierry Coosemans, Salvador García, and Maarten Messagie. Balancing forecast accuracy and switching costs in online optimization of energy management systems. *Expert Systems with Applications*, 298:129305, 2026. doi: 10.1016/j.eswa.2025.129305.
- Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021.
- Isaac Gibbs and Emmanuel Candès. Conformal inference for online prediction with arbitrary distribution shifts. *Journal of Machine Learning Research*, 2024.
- Tilman Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society B: Statistical Methodology*, 69(2): 243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.

- 225 Rakshitha Godahewa, Christoph Bergmeir, Zeynep Erkin Baz, Chengjun Zhu, Zhangdi Song, Salvador
226 García, and Dario Benavides. On forecast stability. *International Journal of Forecasting*, 41(4):
227 1539–1558, 2025. doi: 10.1016/j.ijforecast.2025.01.006.
- 228 Dongjian Hu, Junxi Wu, Shu-Tao Xia, and Changliang Zou. Distribution-informed online conformal
229 prediction. arXiv:2512.07770v2, 2026.
- 230 Ankit Lade, Sai Krishna J., and Indar Kumar. Bias-corrected adaptive conformal inference for
231 multi-horizon time series forecasting. arXiv:2604.13253v1, 2026.
- 232 Ruipu Li and Alexander Rodríguez. Neural conformal control for time series forecasting. In
233 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18439–18447,
234 2025. doi: 10.1609/aaai.v39i17.34029.
- 235 Ruipu Li, Daniel Menacho, and Alexander Rodríguez. Optimization-based online conformal predic-
236 tion for multi-step forecasting. arXiv:2508.13362v3, 2025.
- 237 Yizhou Min, Yizhou Lu, Lanqi Li, Zhen Zhang, and Jiaye Teng. Questioning the coverage-length
238 metric in conformal prediction: When shorter intervals are not better. arXiv:2601.21455v2, 2026.
- 239 F. Pappenberger, H. L. Cloke, A. Persson, and D. Demeritt. HESS Opinions “on forecast
240 (in)consistency in a hydro-meteorological chain: curse or blessing?”. *Hydrology and Earth
241 System Sciences*, 15(7):2391–2400, 2011. doi: 10.5194/hess-15-2391-2011.
- 242 P. Pinson and R. Girard. Evaluating the quality of scenarios of short-term wind power generation.
243 *Applied Energy*, 96:12–20, 2012. doi: 10.1016/j.apenergy.2011.11.004.
- 244 Aleksandr Podkopaev, Darren Xu, and Kuang-Chih Lee. Adaptive conformal inference by betting,
245 2024.
- 246 Kandrika Pritularga and Nikolaos Kourentzes. Forecast congruence: A quantity to align forecasts
247 and inventory decisions. SSRN working paper 4711817, 2024.
- 248 Aaditya Ramdas, Peter Grünwald, Vladimir Vovk, and Glenn Shafer. Game-theoretic statistics and
249 safe anytime-valid inference. *Statistical Science*, 38(4):576–601, 2023. doi: 10.1214/23-STS894.
- 250 Pouria Sarhadi. On the standard performance criteria for applied control design: PID, MPC or
251 machine learning controller? arXiv:2503.14379v3, 2025.
- 252 Glenn Shafer and Vladimir Vovk. *Game-Theoretic Foundations for Probability and Finance*. Wi-
253 ley Series in Probability and Statistics. Wiley, 2019. ISBN 9780470903056. doi: 10.1002/
254 9781118548035.
- 255 Huseyin Tunc, Onur A. Kilic, S. Armagan Tarim, and Burak Eksioglu. A simple approach for
256 assessing the cost of system nervousness. *International Journal of Production Economics*, 141(2):
257 619–625, 2013. doi: 10.1016/j.ijpe.2012.09.022.
- 258 Jente Van Belle, Ruben Crevits, and Wouter Verbeke. Improving forecast stability using deep learning.
259 *International Journal of Forecasting*, 39(3):1333–1350, 2023. doi: 10.1016/j.ijforecast.2022.06.
260 007.
- 261 Jente Van Belle, Honglin Wen, Wouter Verbeke, and Pierre Pinson. Stabilizing distribution-free
262 probabilistic forecasts. arXiv:2605.28531v1, 2026.
- 263 Rahul Vaze. Simultaneous coverage and efficiency guarantee in online conformal prediction.
264 arXiv:2607.26577v1, 2026.
- 265 Vladimir Vovk and Ruodu Wang. E-values: Calibration, combination and applications. *The Annals
266 of Statistics*, 49(3):1736–1754, 2021. doi: 10.1214/20-AOS2020.
- 267 Xiaoqian Wang and Rob J. Hyndman. Online conformal inference for multi-step time series forecast-
268 ing. arXiv:2410.13115v2, 2024.

- 269 Rochelle P. Worsnop, Michael Scheuerer, Thomas M. Hamill, and Julie K. Lundquist. Generating
270 wind power scenarios for probabilistic ramp event prediction using multivariate statistical post-
271 processing. *Wind Energy Science*, 3(1):371–393, 2018. doi: 10.5194/wes-3-371-2018.
- 272 Junxi Wu, Dongjian Hu, Yajie Bao, Shu-Tao Xia, and Changliang Zou. Error-quantified conformal
273 inference for time series. arXiv:2502.00818v2, 2025.
- 274 Marco Zanotti. Analyzing the retraining frequency of global forecasting models: towards more stable
275 forecasting systems. arXiv:2506.05776v3, 2026.
- 276 Mingyi Zheng, Hongyu Jiang, Yizhou Lu, and Jiaye Teng. Smoothing-based conformal prediction
277 for balancing efficiency and interpretability. arXiv:2509.22529v1, 2025.
- 278 Ervin Zsótér, Roberto Buizza, and David Richardson. “Jumpiness” of the ECMWF and Met Office
279 EPS control and ensemble-mean forecasts. *Monthly Weather Review*, 137(11):3823–3836, 2009.
280 doi: 10.1175/2009MWR2960.1.

281 Code and data

282 The simulator, its frozen configuration, the `results/` files behind every figure and table, the gen-
283 erators that read them, and this project’s audit trail across sessions are at [https://\[REPOSITORY](https://[REPOSITORY LINK])
284 [LINK](https://[REPOSITORY LINK]) - INSERTED BY AUTHOR PRIOR TO SUBMISSION].

285 A Supporting measurements

Table 1: Placement A: a one-scalar readout on the *completed* threshold, under one deterministic adversary, on one pass to $T = 10^6$; all eleven arms run are listed. Proposition 2’s bound is 13.2061 at $T = 2 \times 10^5$ and 14.8155 at $T = 10^6$; “ratio” is $\max_t |E_t|$ over it, and “Sat.” the fraction of rounds on which (4) delivers its full $\pm b$. Every number is the *primary* regime’s, defined in the paragraph below, which also gives the $T = 10^4$ excursions and the fitted excursion law.

Readout on the completed threshold	miscov. 2×10^5	miscov. 10^6	$\max_t E_t 10^6$	ratio	Sat. 10^6	travel 10^6
none (control)	0.100035	0.100007	7.80	0.53	0.0000	28,311
partial adj. $w = 0.5$	0.100030	0.100007	7.70	0.52	0.0000	18,113
partial adj. $w = 0.9$	0.100030	0.100007	7.50	0.51	0.0000	4,081
partial adj. $w = 0.99$	0.100040	0.100006	63.00	4.25	0.0007	1,023
partial adj. $w = 0.999$	0.100035	0.100004	623.70	42.10	0.0097	202
growing time constant $p = 0.5$	0.100030	0.100006	12.60	0.85	0.0000	330
growing time constant $p = 0.9$	0.100030	0.100024	53.00	3.58	0.3181	16
running mean	0.099940	0.100010	49.50	3.34	0.3647	9
dead band $\tau = 0.5$	0.100035	0.100005	10.60	0.72	0.0000	2,045
dead band $\tau = 0.9$	0.100060	0.100012	13.20	0.89	0.0038	1,051
dead band $\tau = 1.5$	1.000000	1.000000	900,000	60,747	1.0000	0.5

286 **The primary regime.** Every number in Table 1 is measured under one deterministic adversary,
287 which consumes no randomness and needs no seed. It plays $s_t = \text{clip}(\tilde{q}_t + \varepsilon, \pm b/2)$ with $\varepsilon = 10^{-9}$
288 against the *deployed* threshold, and the indicator is strict. The specification is not a tie rule, and
289 the distinction is measured rather than argued. Scoring an exact-threshold play as a *miscoverage*
290 reproduces every quantity Table 1 reports: 44 of 44 arm-horizon cells identical, all eleven arms at
291 all four horizons. Scoring the same play under the *strict* indicator does not, and it moves every
292 row, because the adversary’s signature move then covers rather than misses. Partial adjustment at
293 $w = 0.999$ excurses 70.80 in place of 623.70, and the failing dead band returns 0.000000 in place of
294 1.000000. That 623.70 turns on the adversary’s ε , not the counting convention. The $\varepsilon \downarrow 0$ limit the
295 harness specifies is the first reading. **No row is convention-free with respect to the second**, which
296 is why the adversary is specified rather than tie-broken. At $\tau = \tau^{*+}$ itself the two readings part: the
297 strict indicator returns 0.000000 where scoring the exact play as a miscoverage returns 1.000000.
298 **Neither of those is coverage.** They are $|E_T| = \alpha T$ and $(1 - \alpha)T$, failures of the same order and
299 opposite sign. Under the strict indicator the exact-threshold play scores $\text{err}_t = \mathbf{1}\{\tilde{q}_t < -b/2\}$,
300 which is also the indicator of the constant legal play $s_t \equiv -b/2$, so that column is Section 3’s *mirror*
301 *adversary* rather than a second convention. It returns between 0.09992 and 0.09997 at $\tau = 0.5, 0.9$,
302 0.95, 0.99 and 0.999, and 0.000000 with $\max_t |E_t| = 10,000$ at $\tau = 1, 1.001, 1.01, 1.05, 1.1$ and

1.5. That locates $\tau^{*-} = 1$ on the same eleven widths Figure 2 uses, closed on the failing side. The convention decides which legal adversary realises the endpoint failure, not whether one occurs.




Figure 2: Realised miscoverage against dead-band width τ at the null scorecaster, $\hat{q} \equiv 0$ with a saturator of level b , under the *specified* adversary, where the upward boundary τ^{*+} of (2) is 1. (a) The full sweep: filled markers run to $T = 10^6$ and the four widest bands, drawn open, to $T = 2 \times 10^5$; blue covers, orange forfeits, and the dashed rule is τ^{*+} . (b) The same setting at $T = 10^5$ over eleven measured widths, placed by $\tau - \tau^{*+}$ on a symlog axis whose linear zone is $|\tau - \tau^{*+}| \leq 10^{-3}$, which magnifies the neighbourhood of τ^{*+} ; visual spacing here is not proportional to true distance in τ . Unlabelled minor ticks mark $\tau = 0.95$ and 1.05 . The grey band spans the widest covering width to the narrowest failing one; Figure 3 gives the other four $(r_t, \hat{q})$ settings. The same eleven widths under the *mirror* adversary, which locates τ^{*-} , are read off in the paragraph above.

Excursions at $T = 10^4$, and where the excursion law holds. Partial adjustment excurses 6.30, 63.00 and 623.70 at $w = 0.9, 0.99$ and 0.999 . The last two are flat in T while Proposition 2's bound grows like $\log T$, which is why the ratio falls with the horizon: at $w = 0.999$ it is 61.08 at $T = 10^4$ and 42.10 at $T = 10^6$. On the control and the four partial-adjustment arms the excursion fits $\max_t |E_t| \approx \max\{0.5 h(T) + 0.9, 0.63/(1 - w)\}$ at every horizon, most loosely at $w = 0.9$, the one arm whose two branches cross inside the horizon range. This is a fit to those five arms, not a derived bound. It does not extend to the other six: the growing time constants, the running mean and the dead bands all excurse further than it predicts.




Figure 3: All 19 dead-band runs, at the five admissible $(r_t, \hat{q})$ settings: rows are settings ordered by τ^* , and the shared horizontal axis is τ . Blue markers cover at $\approx \alpha$ and orange markers return 1.000000; filled markers run to $T = 10^6$ and open markers to $T = 2 \times 10^5$, all under the *specified* adversary. The rule on each row is that setting's upward boundary $\tau^{*+} = \sup_x r_t(x) + \sup_t \hat{q}_t - b/2$, printed at the right, and the tinted band is $\tau \leq \tau^{*+}$. Two rules fall off the axis and are drawn as arrows: $\tau^{*+} = 0$ at $\hat{q} \equiv -b/2$, and τ^{*+} unbounded under Angelopoulos et al.'s tangent integrator. The mirror boundary τ^{*-} of (2) is not drawn here and differs from τ^{*+} at the two non-null scorecasters.

Realised miscoverage at the covering settings of Figure 3. At $\hat{q} \equiv +b/2$ the $\tau = 1.5$ band returns 0.100010 with $\max_t |E_t| = 11.20$, inside the bound. Under Angelopoulos et al.'s tangent integrator the same width returns 0.100001. Raising the saturator's level to $4b$ leaves coverage intact and cuts the $w = 0.999$ excursion from 623.70 to 120.60. That is the first term of τ^* moving the edge, not the readout changing behaviour.

318 **What survives of the measurement, and what its status now is.** Outside the admissible window
319 the transition is a step. Past τ^{*+} the threshold sits permanently below the reachable score range and
320 $|E_t| = (1 - \alpha)t$ grows linearly, 60,747 times the bound at $T = 10^6$; the wider bands $\tau = 2.0, 2.5,$
321 3.0 and 5.0 fail too at $T = 2 \times 10^5$ (Figure 2). **Here $\sup_x r_t(x) = b$ and $\hat{q} \equiv 0$, so $\tau^{*+} = \tau^{*-} = b/2$,**
322 **one point of the admissible class, and both terms move it.** The legal $\hat{q} \equiv +b/2$ carries the $\tau = 1.5$
323 band inside the bound against the specified adversary, and against the mirror one no width covers
324 there, the unsmoothed control included. Condition (4) bounds $\sup_x r_t(x)$ only from below, so raising
325 it costs the failure direction nothing; it moves the edge too (Figure 3) and shrinks the excursion,
326 from 623.70 to 20.10 under the tangent integrator Angelopoulos et al. [2023] report using in all their
327 experiments, both at $T = 10^6$. **The method’s own integrator therefore sits inside the set at every**
328 **width run here. The endpoint is where the two boundaries part.** The upper one is closed and the
329 lower one is open, which the null scorecaster cannot show because τ^{*+} and τ^{*-} coincide there: four
330 saturators with $\tau^{*+} < \tau^{*-}$ cover at $\tau = \tau^{*+}$ exactly under both adversaries, and at $\hat{q} \equiv +0.4$, where
331 $\tau^{*-} = 0.6 < 1.4 = \tau^{*+}$, the width 0.599 covers under both while 0.6 fails under the mirror. At the
332 baseline the two coincide at 1, so the primary regime covers there and the mirror adversary does
333 not; Table 1 and the primary-regime paragraph above print both readings. **The failing arms have**
334 **Sat. = 1.0000.** Condition (4) delivers its full $\pm b$ on every round, and what fails is Proposition 2’s
335 other hypothesis: the threshold closing the loop *is* the integrator’s output.

Table 2: Legal $(r_t, \hat{q})$ configurations on which the printed τ^* over-certifies, each under two legal adversaries. “Proved window” is where (2) certifies retention against every legal adversary; between it and the failure boundaries $\tau^{*\pm}$ neither direction is claimed, and the rows built on a supremum the dynamics cannot reach or never attain sit there. Here $\emptyset$ marks a configuration for which no width is certified, $\tau = 0$ included. $T = 10^5$ except the two power-of-two rows, at $T = 10^6$. The last row carries no dead band at all.

$(r_t, \hat{q})$ configuration	τ	printed τ^*	proved window	$s_t = \hat{q}_t + \varepsilon$		$s_t \equiv -b/2$	
				miscov.	$\max_t E_t $	miscov.	$\max_t E_t $
$\hat{q} \equiv +0.4$, level b	0.599	1.4	$\tau < 0.6$	0.100010	7.20	0.099890	11.60
$\hat{q} \equiv +0.4$, level b	0.8	1.4	$\tau < 0.6$	0.100050	8.70	0.000000	10,000
$\sup_x r_t = 10b$, off the reach	1.001	19	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t$ unbounded, off the reach	1.5	∞	$\tau < 1$	1.000000	90,000	0.000000	10,000
$A^+ = 4b, A^- = -b$	1.5	7	$\tau < 1$	0.100020	4.30	0.000000	10,000
level $4b, \hat{q} \equiv 0$	6.999	7	$\tau < 7$	0.100070	12.20	0.100050	11.60
level $4b, \hat{q} \equiv 0$	7.0	7	$\tau < 7$	0.100090	12.20	0.000000	10,000
$\sup_x r_t = 2b$, unattained	2.99	3	$\tau < 1$	0.100320	211.10	0.098010	210.50
$\sup_x r_t = 2b$, unattained	3.0	3	$\tau < 1$	1.000000	90,000	0.000000	10,000
$\sup_x r_t = 2b$, attained	3.0	3	$\tau < 3$	0.100000	12.20	0.000000	10,000
$\hat{q}_t = b/2$ on powers of 2	1.0	2	$\emptyset$	0.100009	14.30	0.000000	100,000
$\hat{q}_t = b/2$ on powers of 2	1.9	2	$\emptyset$	0.110960	84,745.70	0.000000	100,000
$\hat{q} \equiv +b/2$, no band	0	2	$\emptyset$	0.100000	0.90	0.000000	10,000

336 B The four vocabularies

Table 3: Four literatures name the movement of a sequentially updated forecast in four vocabularies: *abrupt threshold changes*; *forecast (in)stability* (MASC, MAC/SDC, MQC/SMQC, congruence, nervousness); *jumpiness*, *convergence index*, *(in)consistency*; and *switching cost*. The row gives what each contributes.

Online conformal	Forecast stability	Verification	Switching-cost online learning
Distribution-free long-run coverage for any bounded scorecaster as post-processing [Angelopoulos et al., 2023, Thm 1, Prop. 2]	Movement functional; readout as post-processing [Godaheva et al., 2025, Eq. 3]; priced decision [Genov et al., 2026, §4.2], [Van Belle et al., 2026]	Matched-level design, real producers (2012) [Pinson and Girard, 2012, Worsnop et al., 2018]	Worst-case theory for the penalised problem [Andrew et al., 2013, Thm 2]

337 C Where this sits

338 Four vocabularies name one quantity; Table 3 (App. B) names all four. **What is conceded, by name.**
339 The placement is prior work: Angelopoulos et al. [2023, Thm. 1, Prop. 2] state coverage, with a
340 rate, for any admissible integrator and any scorecaster, and run a smoothing-family model in that
341 slot themselves; this paper proves that admissible set tight. Dupuy et al. [2026, Appendix A] give
342 the generic argument, and Duerst et al. [2024] adopt that scorecaster. The derivation is prior work
343 too, and this concession matters most. Li et al. [2025, Corollary 2] prove the retention statement in
344 Conformal PID’s own notation, with Wang and Hyndman [2024] an independent second instance.
345 Section 2 concedes the smoother object, the metric arithmetic and both readout maps. Zheng et al.
346 [2025], Wu et al. [2025], Chen et al. [2023], Lade et al. [2026] each place a smoother inside their
347 own validity argument. One vocabulary out, this is integrator anti-windup with a feedforward path:
348 free-until-the-bound-binds is its defining specification, not a finding [Galeani et al., 2009]. Nearest
349 in the other direction, Gao et al. [2025] run the same recursion but gate the *model* update on the
350 conformal score. Their one non-standard assumption is that $\sum_{t \leq T} q_t \leq O(\alpha h(T))$ for a sublinear
351 h , supported with a plot rather than derived. Verification’s matched-level design tests real producers
352 from 2012, with a control stronger than the pair: the whole marginal [Pinson and Girard, 2012,
353 Worsnop et al., 2018]. Switching-cost online learning adds a Thm. 7 trade-off in which sublinear
354 regret is bought by growing θ with T , which grows the ratio [Andrew et al., 2013, Thm. 7].

355 **The accumulator, read as a bet.** A Skeptic who each round pays α for the payoff err_t has cumulative
356 net gain exactly E_t , so this paper’s accumulator is a game-theoretic capital process at unit stake
357 [Shafer and Vovk, 2019, Ramdas et al., 2023]. It is not an e-value and not an e-process: it is signed,
358 and Proposition 2 bounds it pathwise by feedback rather than under a null by Ville’s inequality [Vovk
359 and Wang, 2021]. What is formal is the standard transform. Under the null that $\text{err}_t \mid \mathcal{F}_{t-1}$ is
360 Bernoulli(α), $M_t = \prod_{i \leq t} (1 + \lambda(\text{err}_i - \alpha))$ is a test martingale for each $\lambda \in [-1/(1 - \alpha), 1/\alpha]$,
361 with $\log M_t = \lambda E_t + O(\lambda^2 t)$. Section 3’s boundary is stated in $|E_T|$ and proved in both directions,
362 so it transfers as a dichotomy rather than as a measured rate: inside the window $|E_T| \leq c h(T) + 1$
363 and $M_T \rightarrow 0$ exponentially at every $\lambda \neq 0$; past it $|E_T|$ is linear in T and some fixed λ grows M_T
364 exponentially, so an anytime-valid test of the deployed intervals’ calibration rejects at any fixed level
365 within $O(1)$ rounds. Betting enters online conformal inference elsewhere as an optimiser rather than
366 as evidence [Podkopaev et al., 2024]. We construct no e-value here and claim none.